

# MPI Reference

## General:

### Initialize MPI:

```
int MPI_Init(int *argc,
             char ***argv)
```

argc and argv come from the main function input argument, can use MPI\_Init(NULL, NULL) if necessary.

### Cleanup:

```
int MPI_Finalize()
```

All processes need to call this before exiting.

### How many processes in the communicator?:

```
int MPI_Comm_size(MPI_Comm comm,
                  int *size)
```

### Which rank in communicator am I?:

```
int MPI_Comm_rank(MPI_Comm comm,
                  int *rank)
```

### Check wall clock time:

```
double MPI_Wtime()
```

## Datatypes:

|            |          |
|------------|----------|
| MPI_INT    | int      |
| MPI_LONG   | long int |
| MPI_FLOAT  | float    |
| MPI_DOUBLE | double   |
| MPI_CHAR   | char     |

## Communicators:

MPI\_COMM\_WORLD all processes available to the program

We can define own communicators (TBD).

## Blocking Point-to-point communication:

Blocking means once the code enters the command, it waits until it is completed before moving on with execution of the rest of the program.

### Send a message to one process:

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype,
             int dest, int tag, MPI_Comm comm)
```

Sends count elements of type datatype from the memory location buf to process dest in communicator comm. The tag needs to be matched by the MPI\_Recv command.

### Receive a message from one process:

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,
             int source, int tag, MPI_Comm comm,
             MPI_Status *status)
```

Receives count elements of type datatype to the memory location buf from process source in communicator comm. The tag and source need to match tag and dest from the corresponding MPI\_Send command, unless MPI\_ANY\_TAG and/or MPI\_ANY\_SOURCE wildcards are used instead. If the MPI\_Recv does not match MPI\_Send, the process will hang indefinitely. The status structure holds information about the number of elements received, the tag used and the sender rank. To ignore status, use MPI\_STATUS\_IGNORE.

### Combined Send-Receive:

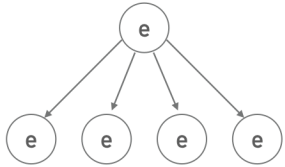
```
int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype
                sendtype, int dest, int sendtag,
                void *recvbuf, int recvcount, MPI_Datatype
                recvtype, int source, int recvtag,
                MPI_Comm comm, MPI_Status *status)
```

Send sendcount elements of type sendtype stored in sendbuf to process dest with tag sendtag, at the same time receive recvcount elements of recvtype into recvbuf from process source with tag recvtag. All happening within communicator comm.

It is possible to reuse the same buffer for sendbuf and recvbuf - see MPI\_Sendrecv\_replace.

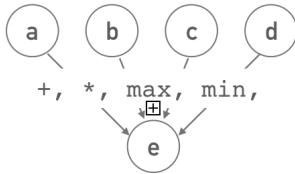
## Collective Communication #1

All processes in a communicator need to call the collective communication command.



### Send message to all processes:

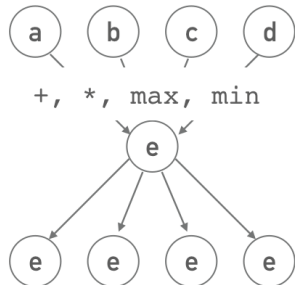
`int MPI_Bcast(void *buf, int count, MPI_Datatype datatype, int root, MPI_Comm comm)`  
Content of buf (a total of count elements of type datatype) is sent from root to all processes in comm. All processes have the copy of data in buf.



### Collect data from all processes and perform an operation to combine it:

`int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op operation, int root, MPI_Comm comm)`

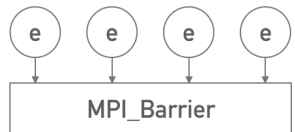
Each process in comm sends the content of sendbuf (count elements of type datatype). An operation is performed on the data, and the result is stored in recvbuf on process root. Some available reduction operations: `MPI_SUM`, `MPI_PROD`, `MPI_MAX`, `MPI_MIN`



### Collect data from all processes, perform reduction operation and store result on all processes:

`int MPI_Allreduce(const void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)`

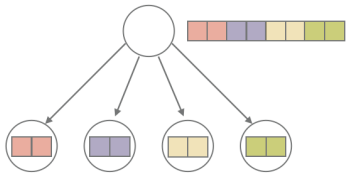
Each process in comm sends the content of sendbuf (count elements of type datatype). An operation is performed on the data, and the result is stored in recvbuf on all process in comm.



### Wait for all processes to reach this point in the program

`int MPI_Barrier(MPI_Comm comm)`

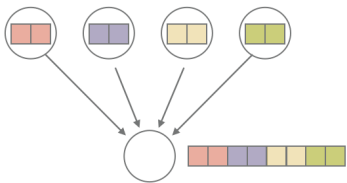
## Collective Communication #2



**Send equal part of a buffer to all processes:**

```
int MPI_Scatter(void *sendbuf, int sendcount, MPI Datatype sendtype,  
                void *recvbuf, int recvcount, MPI Datatype recvtype,  
                int root, MPI Comm comm )
```

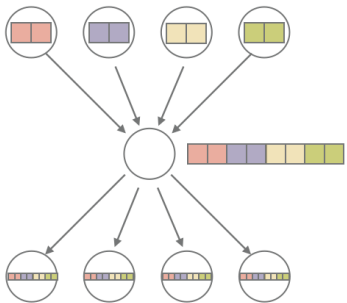
Content of sendbuf on process root is split into equal parts of size sendcount elements of type sendtype and each chunk is sent to a different process in comm, where it is stored in \*recvbuf. Typically recvcount = sendcount, and recvtype = sendtype. Root process also gets its chunk of data send to self.



**Collect data from all processes into a single array on one process**

```
int MPI_Gather(void *sendbuf, int sendcount, MPI Datatype sendtype,  
               void *recvbuf, int recvcount, MPI Datatype recvtype,  
               int root, MPI Comm comm )
```

Small chunks (sendcount elements of type sendtype) of data stored in sendbuf on all processes in comm are collected into one array recvbuf on root process. Typically recvcount = sendcount, and recvtype = sendtype. Unlike MPI\_Reduce, no operation is performed, but the data from each process is appended to an array, which is nproc\*recvcount\*sizeof(recvtype) long.



**Collect data from all processes into a single array copied on all processes**

```
int MPI_Allgather(void *sendbuf, int sendcount, MPI Datatype sendtype,  
                  void *recvbuf, int recvcount, MPI Datatype recvtype,  
                  MPI Comm comm )
```

Small chunks (sendcount elements of type sendtype) of data stored in sendbuf on all processes in comm are collected into one array recvbuf, which is stored on all processes - each process has exact copy of recvbuf. Typically recvcount = sendcount, and recvtype = sendtype.