




ME 471/571

Week 13 - Shared memory and thread synchronization

THREADS SYNCHRONIZATION


Consider a computation of dot product:



```
void dotProduct(float *A, float *B, float *C, const int N)
for (int idx = 0; idx < N; idx++)
    *C += A[idx] * B[idx];
```

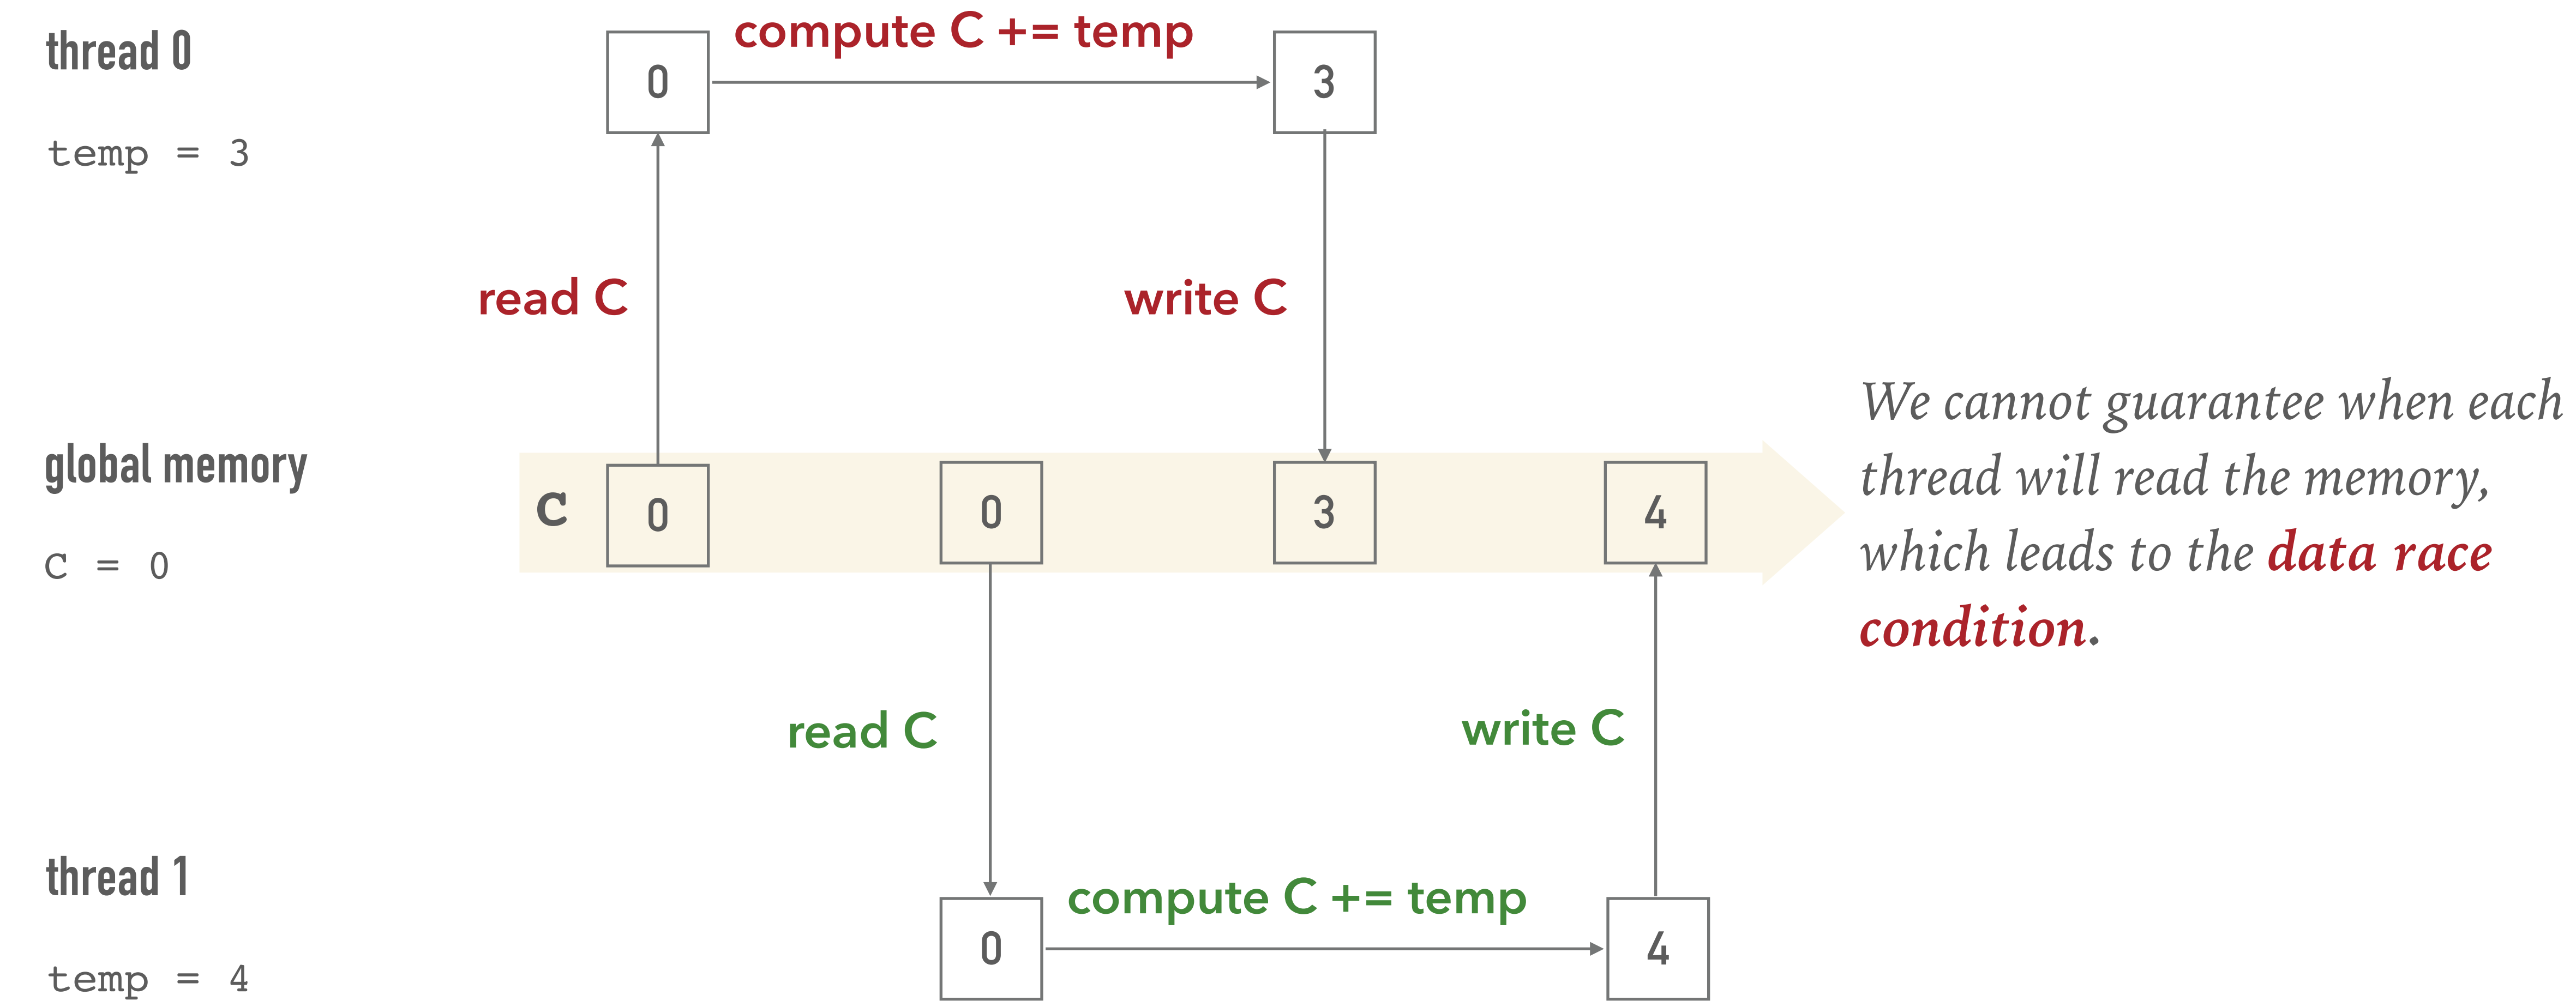
THREADS SYNCHRONIZATION

Parallel threads have no problem with the multiplication part:

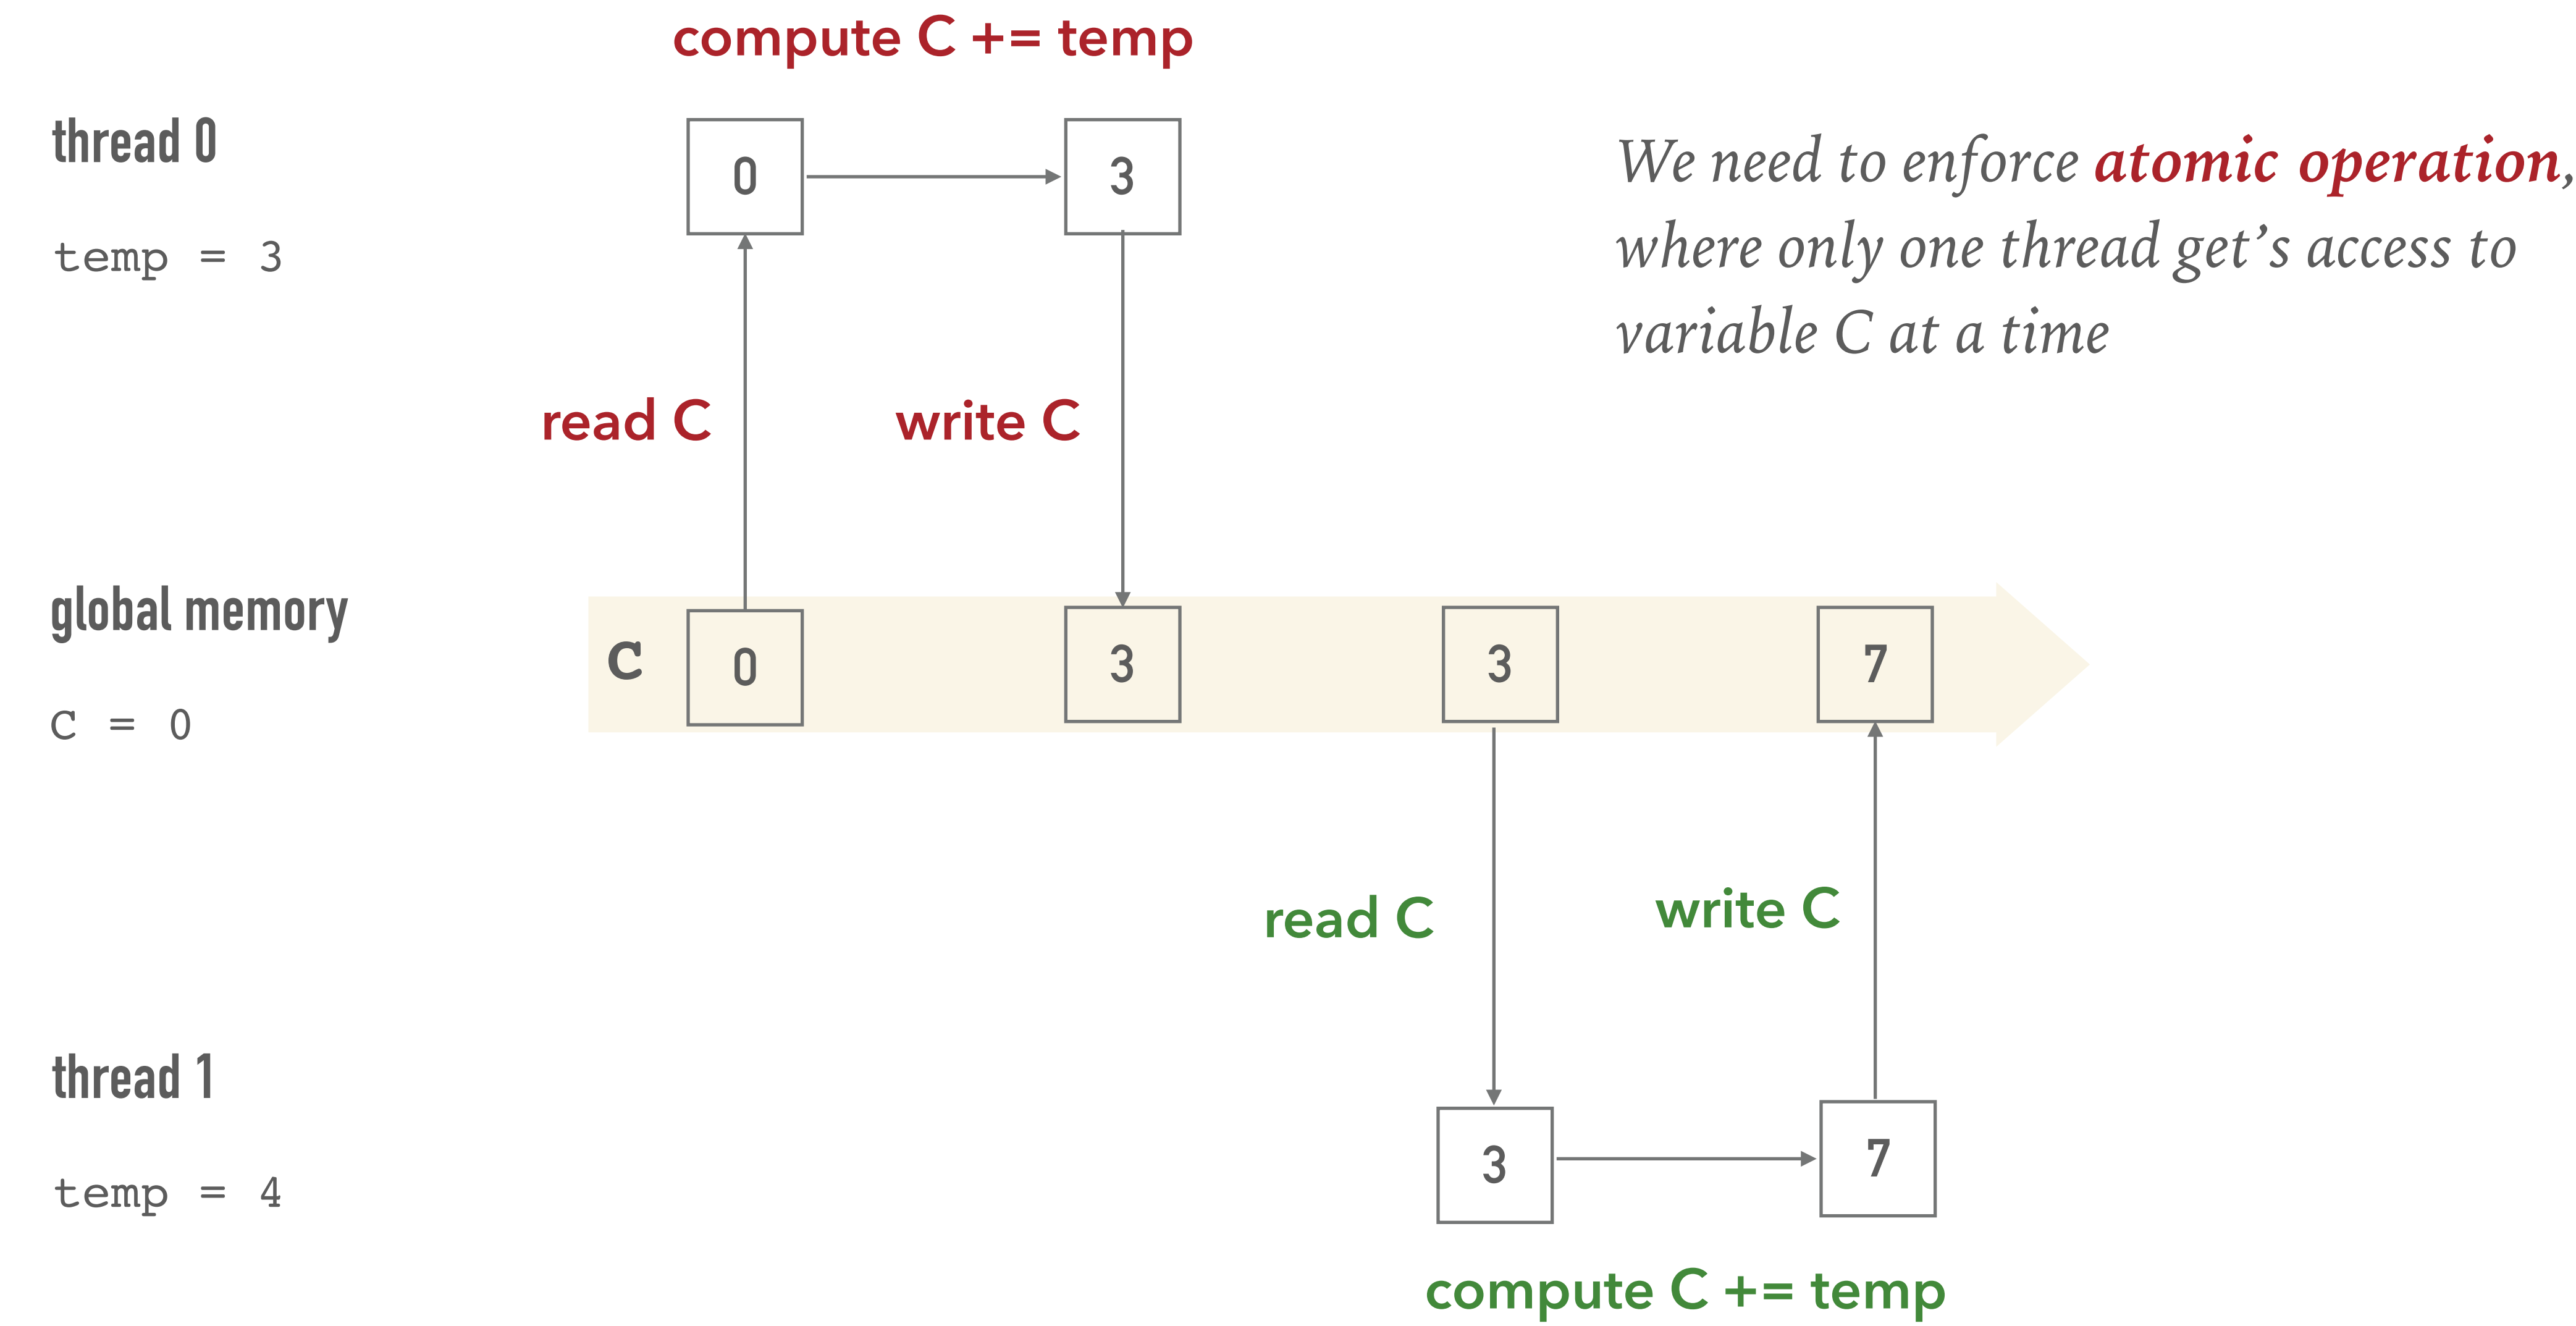


```
__global__ void dot( float *a, float *b, float *c, int N ) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    // Each thread computes a pairwise product  
    float temp = a[idx] * b[idx];  
    c += temp;  
}
```

DATA RACE CONDITION



DATA RACE CONDITION



ATOMIC OPERATIONS

Atomic operations ensure that only one thread can access and modify a memory location. Other threads need to wait until atomic operation is completed.

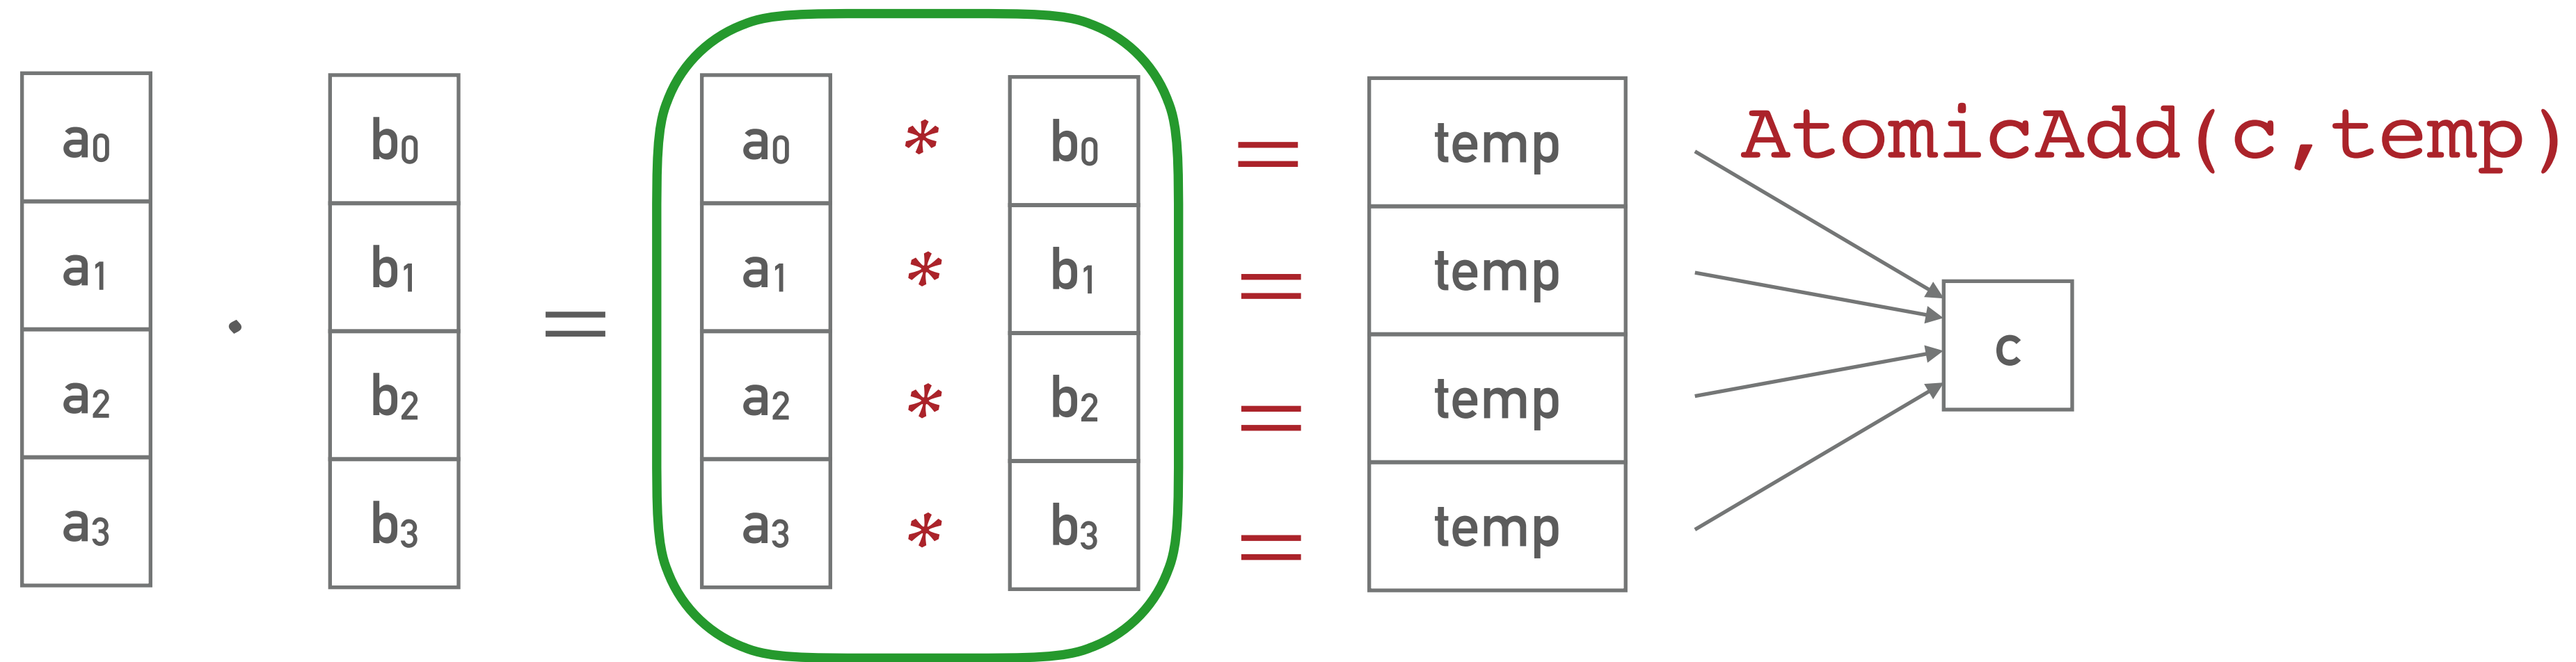
- | | | | |
|-------------|--------------------|--------------|-------------------------|
| ➤ atomicAdd | <i>addition</i> | ➤ atomicInc | <i>increment</i> |
| ➤ atomicSub | <i>subtraction</i> | ➤ atomicDec | <i>decrement</i> |
| ➤ atomicMin | <i>minimum</i> | ➤ atomicExch | <i>exchange</i> |
| ➤ atomicMax | <i>maximum</i> | ➤ atomicCAS | <i>compare and swap</i> |

atomicAdd_block *atomic with respect to block only*

atomicAdd_system *atomic with respect to GPU and CPU (for unified memory)*

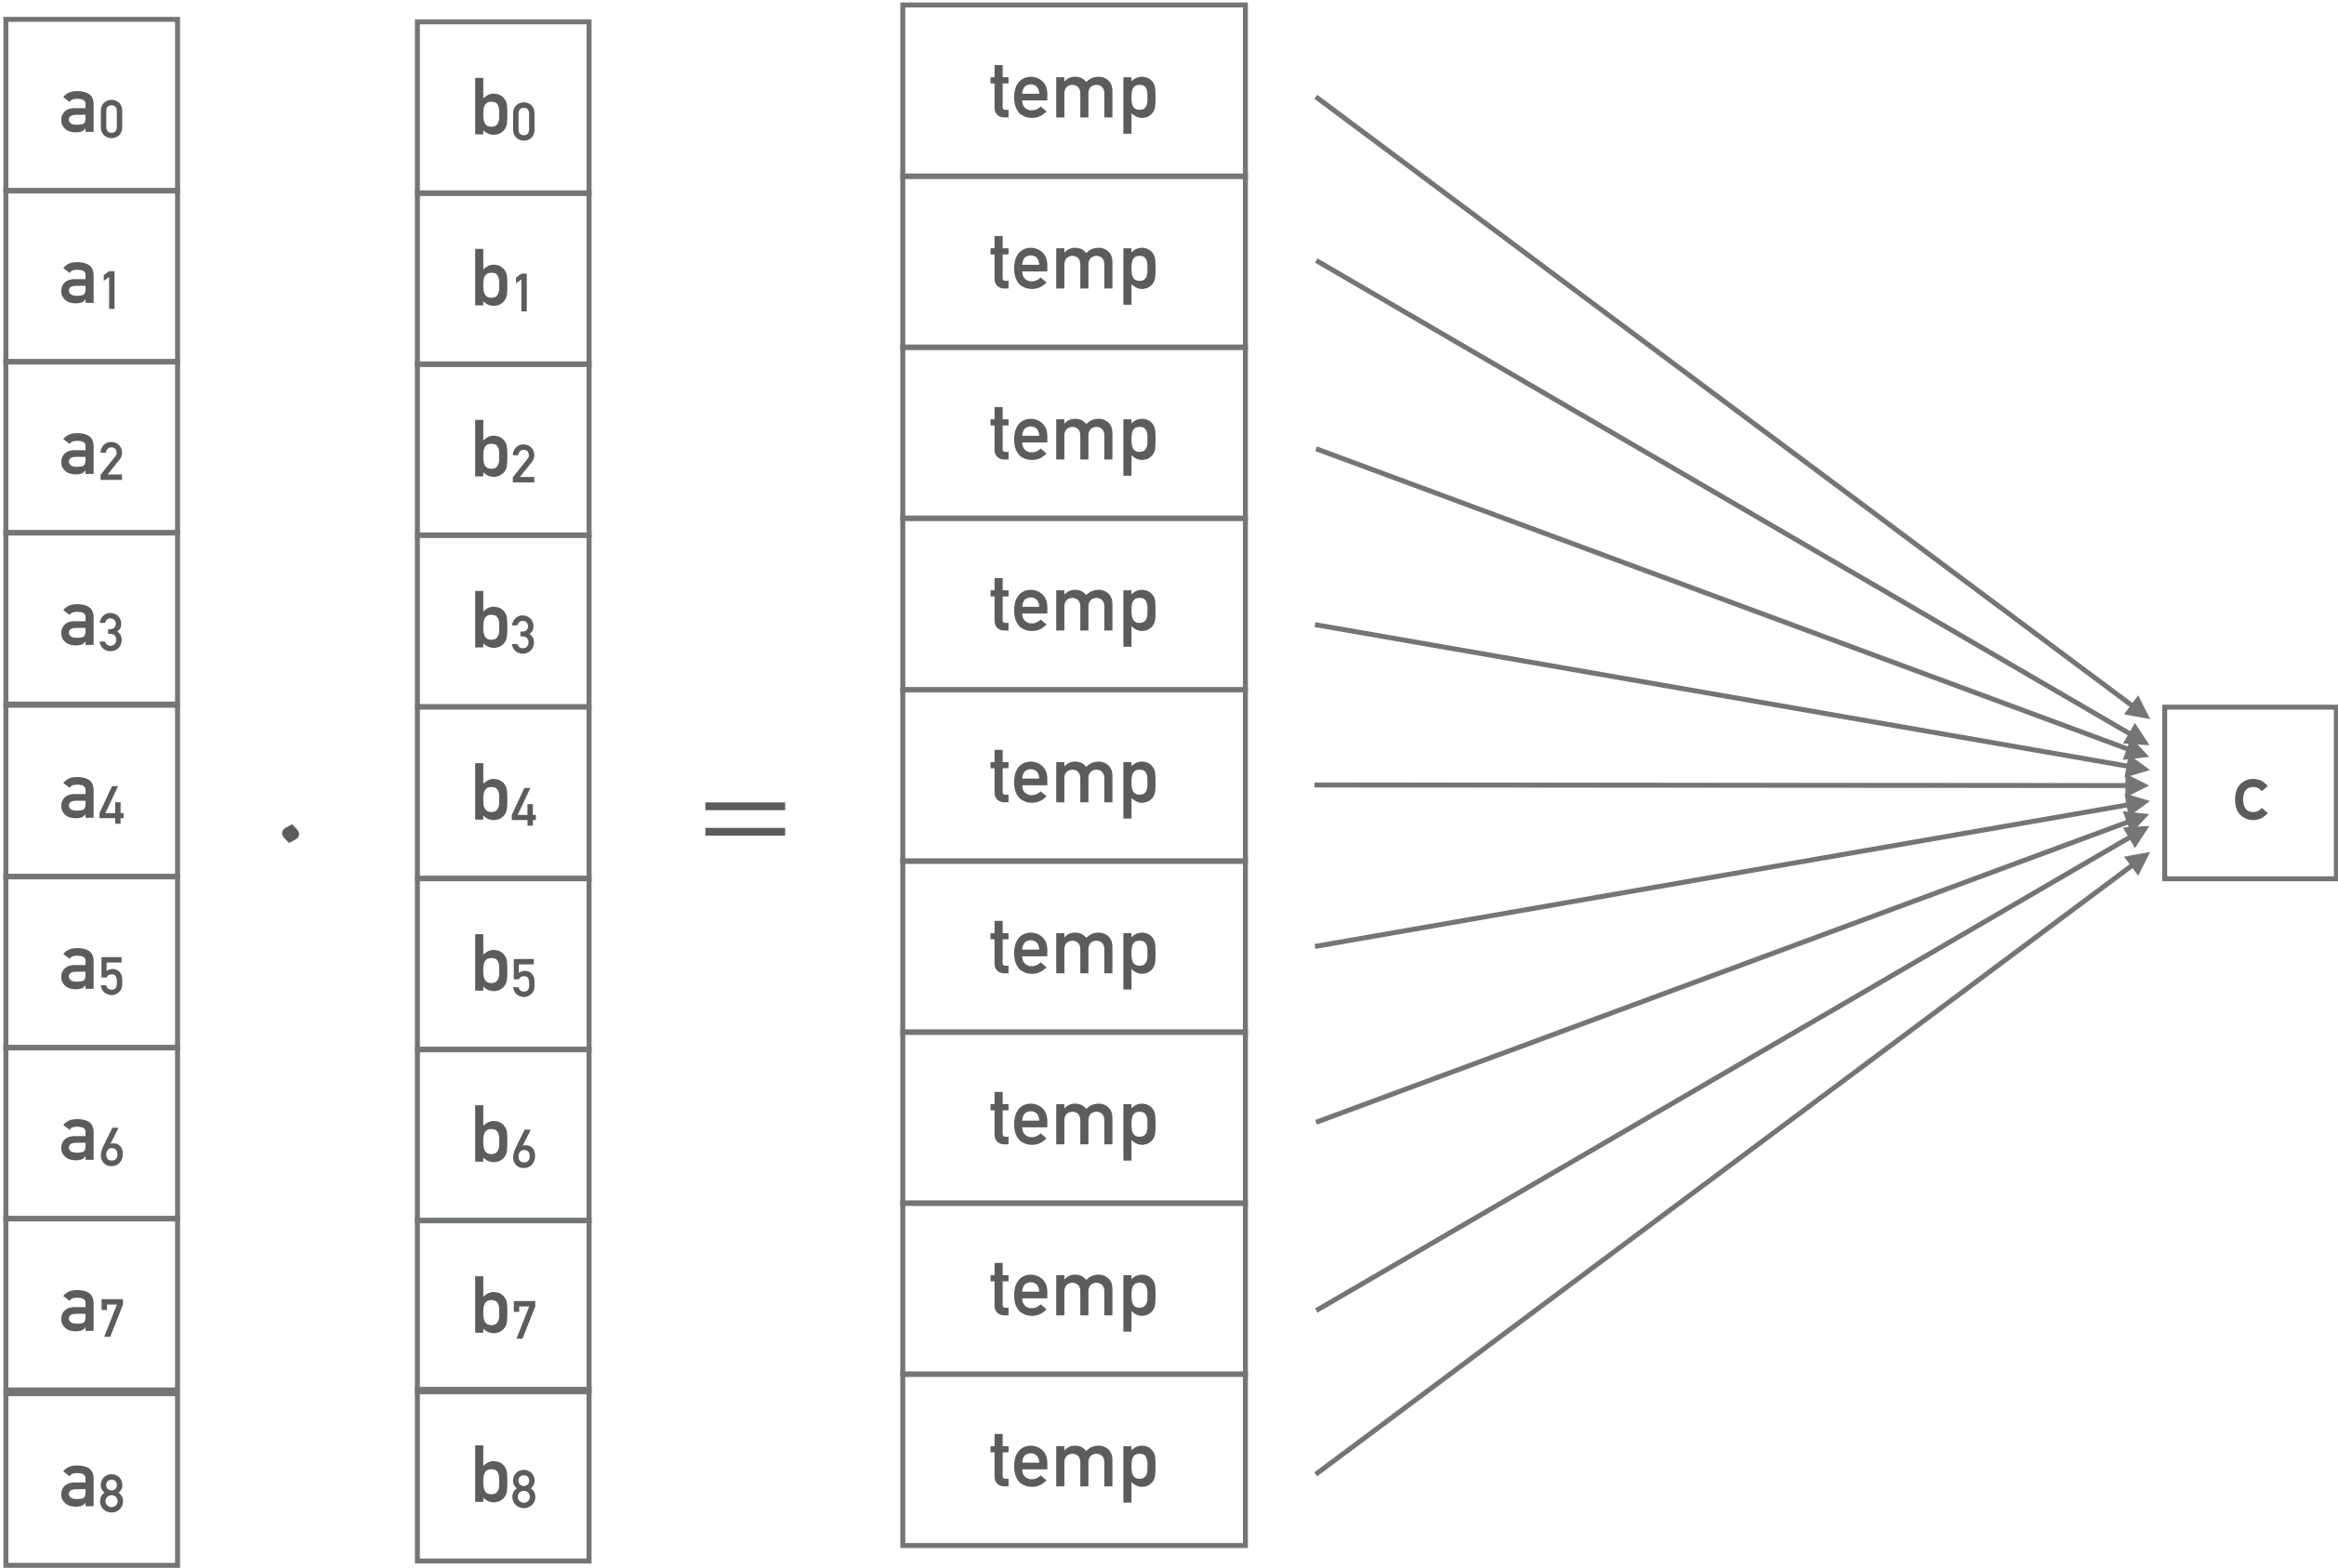
THREADS SYNCHRONIZATION

Parallel threads have no problem with the multiplication part:

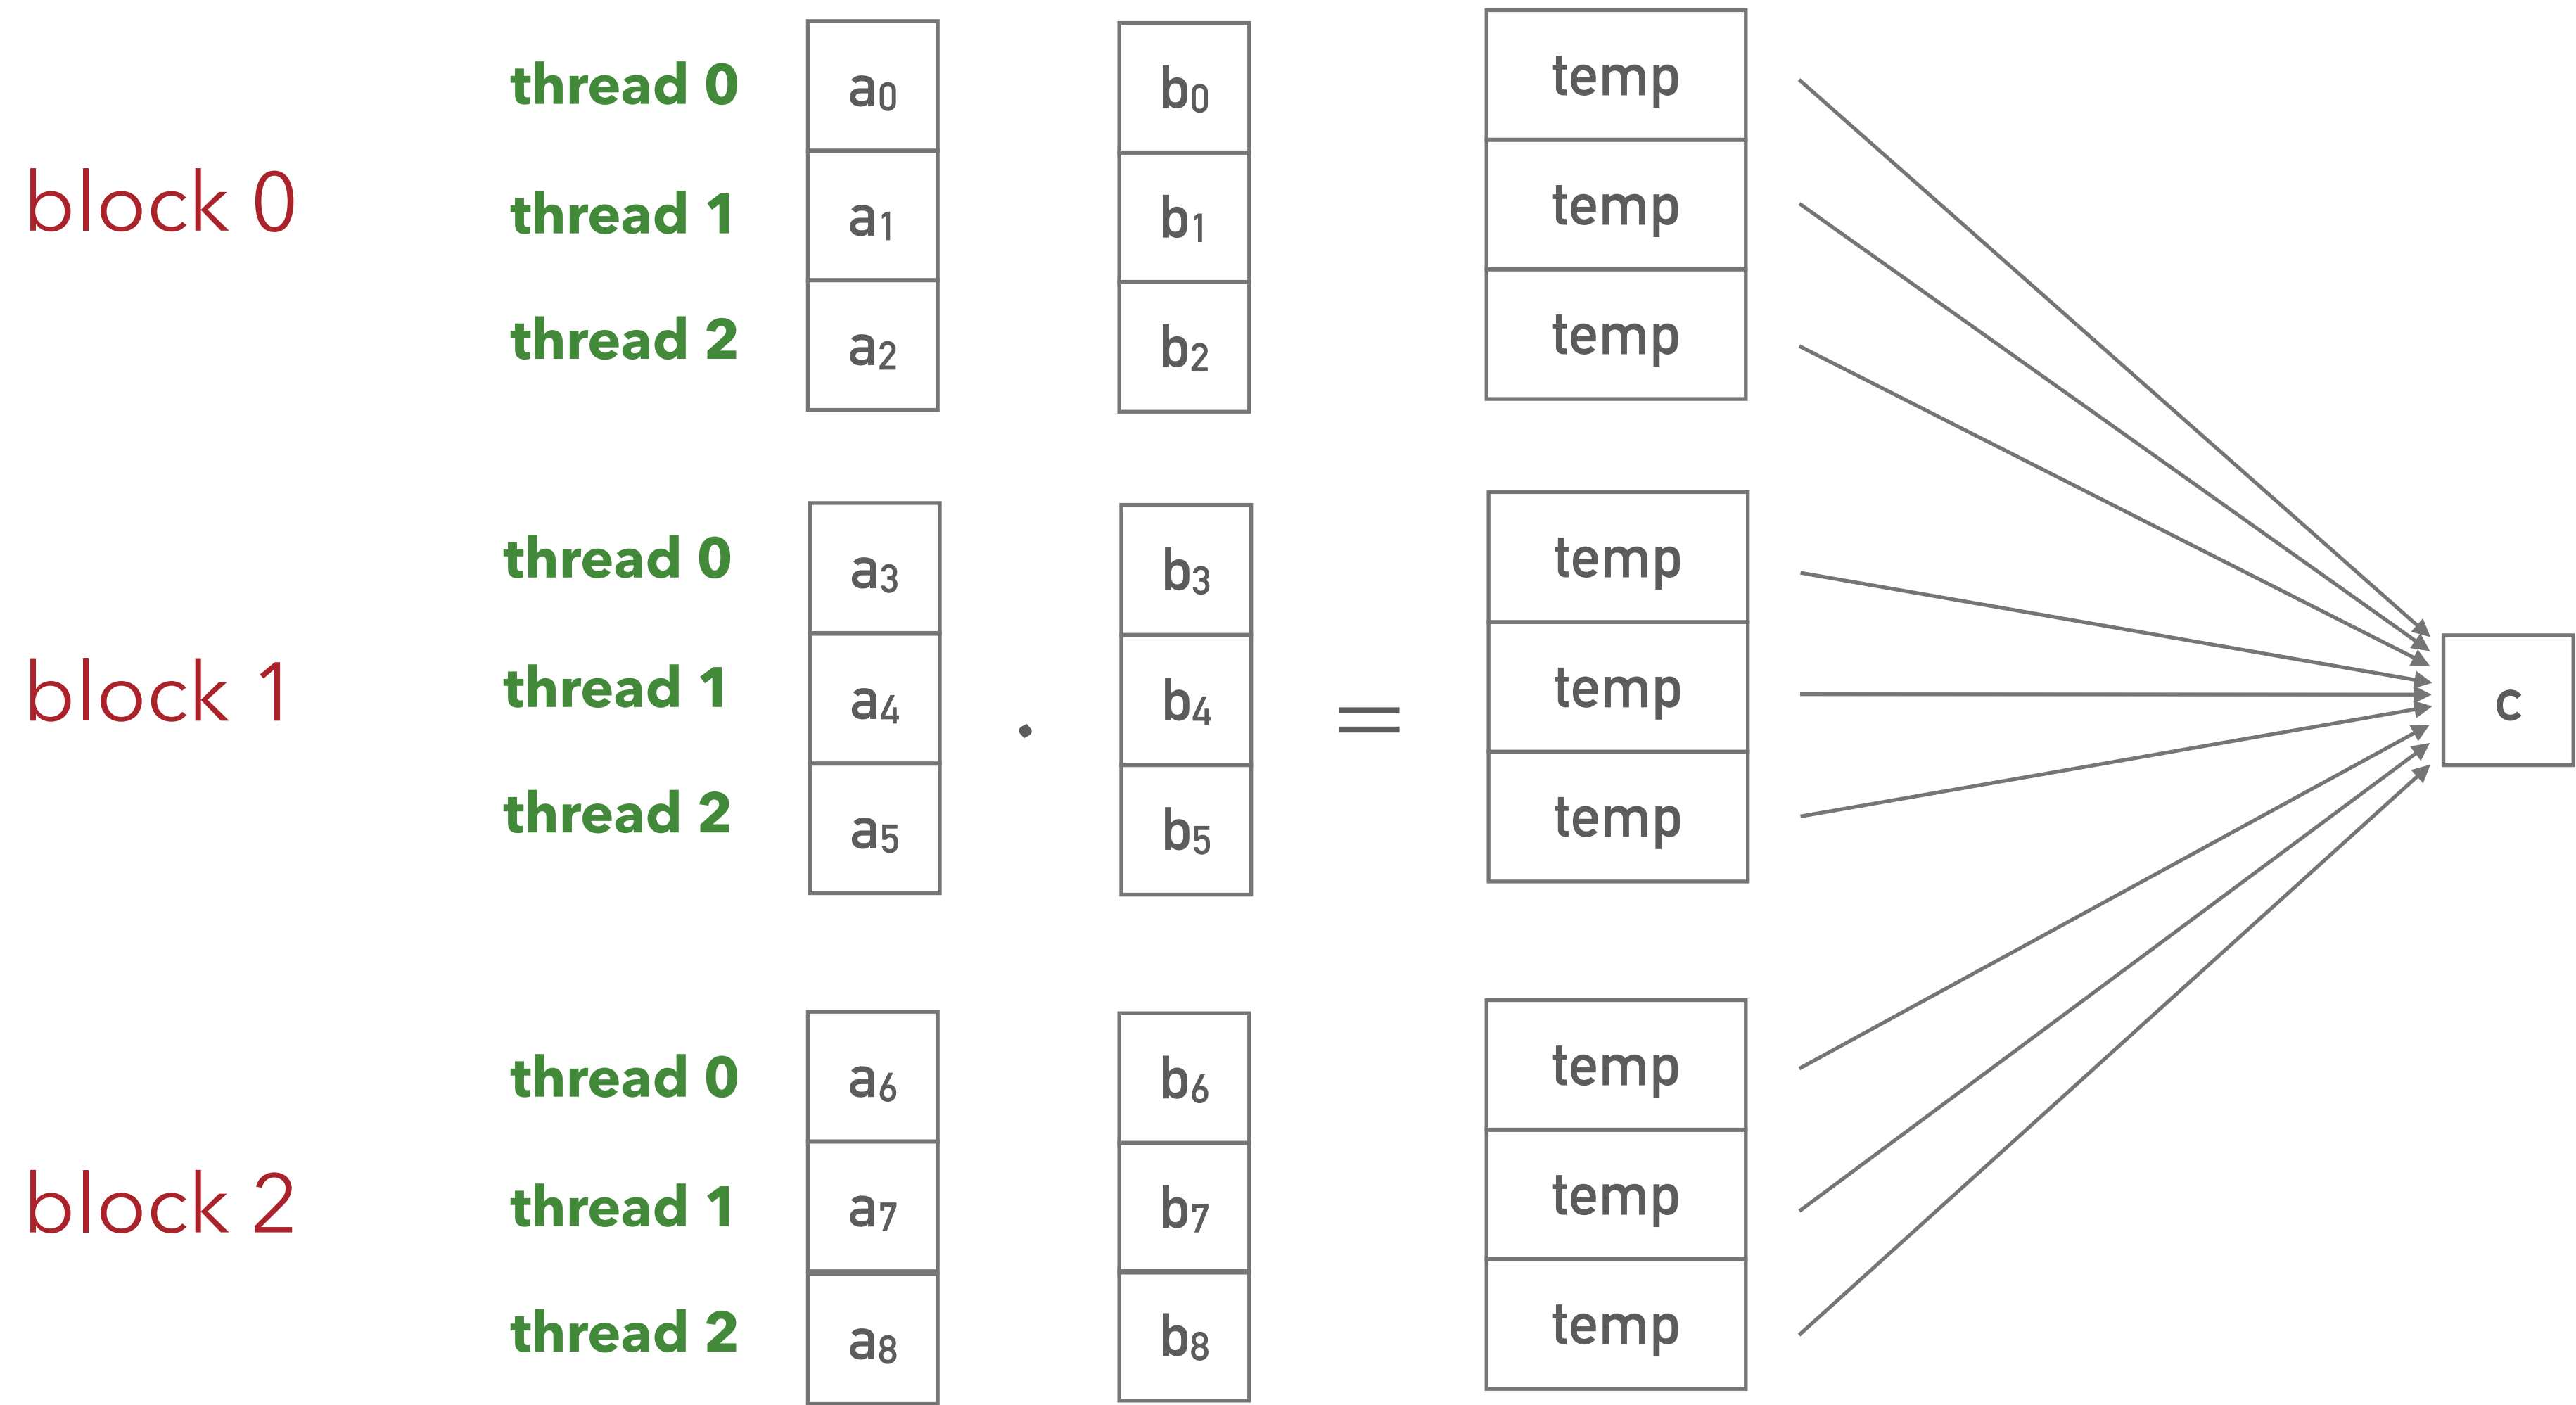


```
__global__ void dot( float *a, float *b, float *c, int N ) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    // Each thread computes a pairwise product  
    float temp += a[idx] * b[idx];  
    AtomicAdd(c, temp);  
}
```

BLOCKS AND THREADS

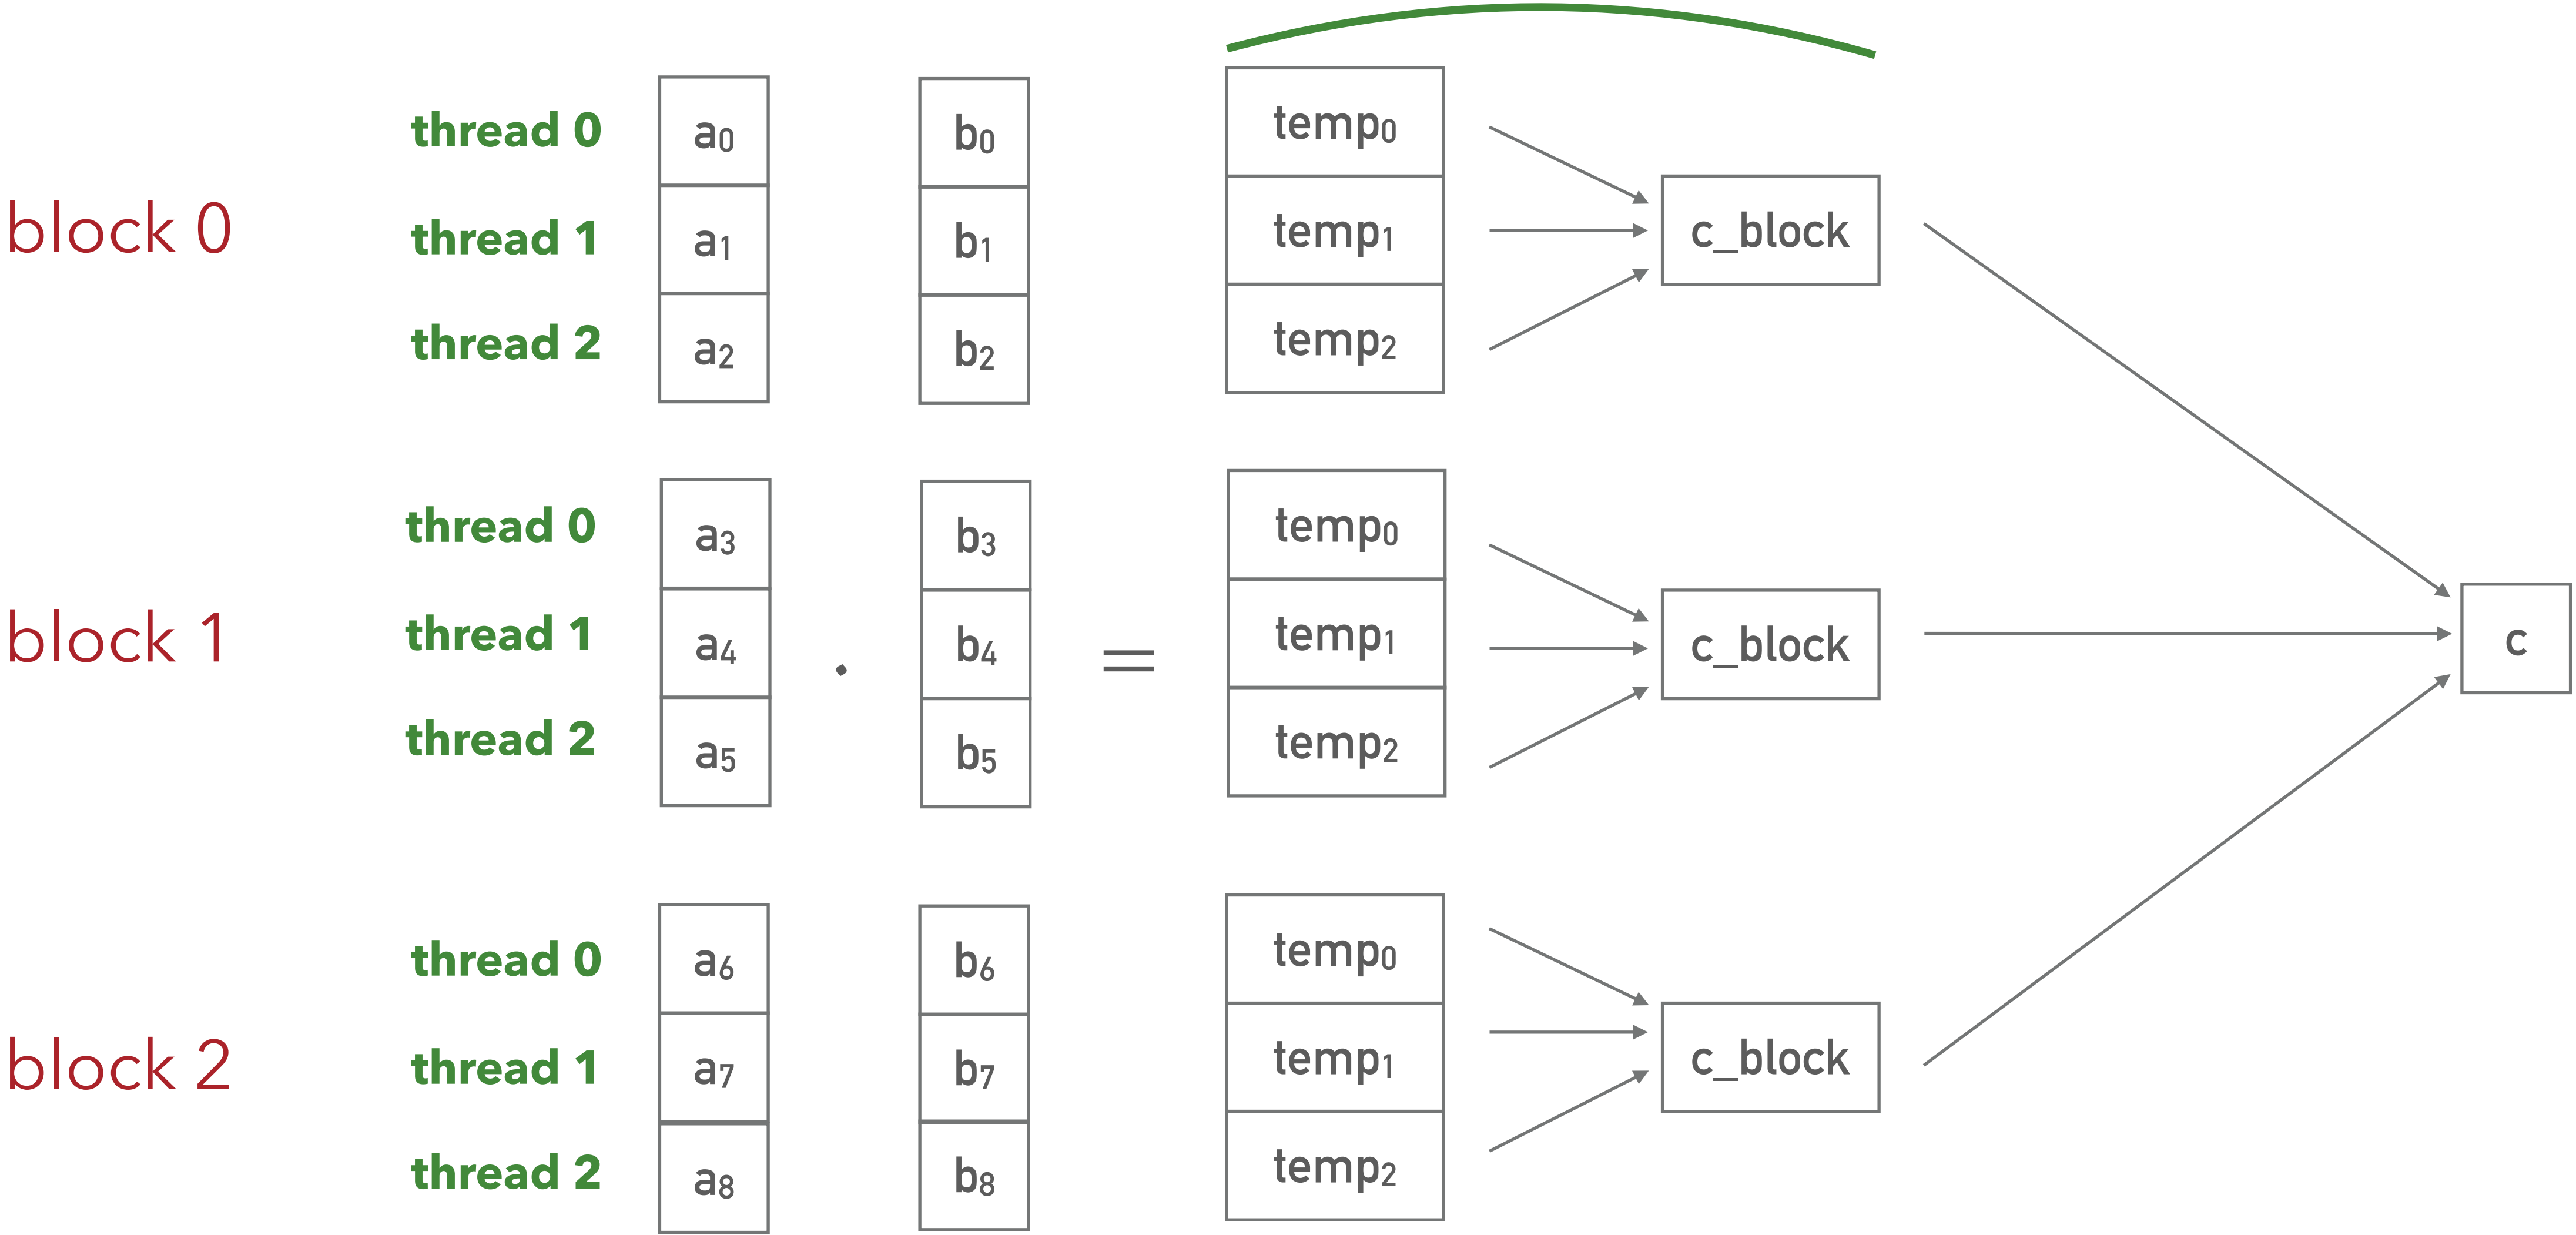


BLOCKS AND THREADS



BLOCKS AND THREADS

shared memory within a block



THREADS SYNCHRONIZATION

Shared memory:

- is shared among the threads, but private to each block
- is extremely fast (think cache, but shared among threads)

```
__global__ void dot( int *a, int *b, int *c, int N ) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
  
    // Each thread computes a pairwise product  
    __shared__ int temp[THREADS_PER_BLOCK];  
    temp[threadIdx.x] = a[idx] * b[idx];  
}
```

THREADS SYNCHRONIZATION

```
__global__ void dot( int *a, int *b, int *c, int N ) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;


    // Each thread computes a pairwise product
    __shared__ int temp[TREADS_PER_BLOCK];
    temp[threadIdx.x] = a[idx] * b[idx];

    __syncthreads();

    // Thread 0 sums up the pairwise products
    if(threadIdx.x == 0) {
        float block_c = 0;
        for (int i = 0; i<N; i++)
            block_c += temp[i];
        atomicAdd(c,block_c);
    }
}
```


FINITE DIFFERENCE ON GPUS

$$\frac{\partial u}{\partial x} \approx \frac{u_{i+1} - u_{i-1}}{2\Delta x}$$



FINITE DIFFERENCE ON GPUS


$$\frac{\partial u}{\partial x} \approx \frac{u_{i+1} - u_{i-1}}{2\Delta x}$$



this is fine, since we have global memory


FINITE DIFFERENCE ON GPUS

$$\frac{\partial u}{\partial x} \approx \frac{u_{i+1} - u_{i-1}}{2\Delta x}$$




we need to do boundary points using one-sided difference, as usual


FINITE DIFFERENCE ON GPUS – SHARED MEMORY




FINITE DIFFERENCE ON GPUS – SHARED MEMORY




FINITE DIFFERENCE ON GPUS – SHARED MEMORY



FINITE DIFFERENCE ON GPUS – SHARED MEMORY



FINITE DIFFERENCE ON GPUS – SHARED MEMORY – GHOSTS



FINITE DIFFERENCE ON GPUS – SHARED MEMORY – GHOSTS

