ME 471/571

Week 3 - Collective Communication

MESSAGE PASSING MODEL

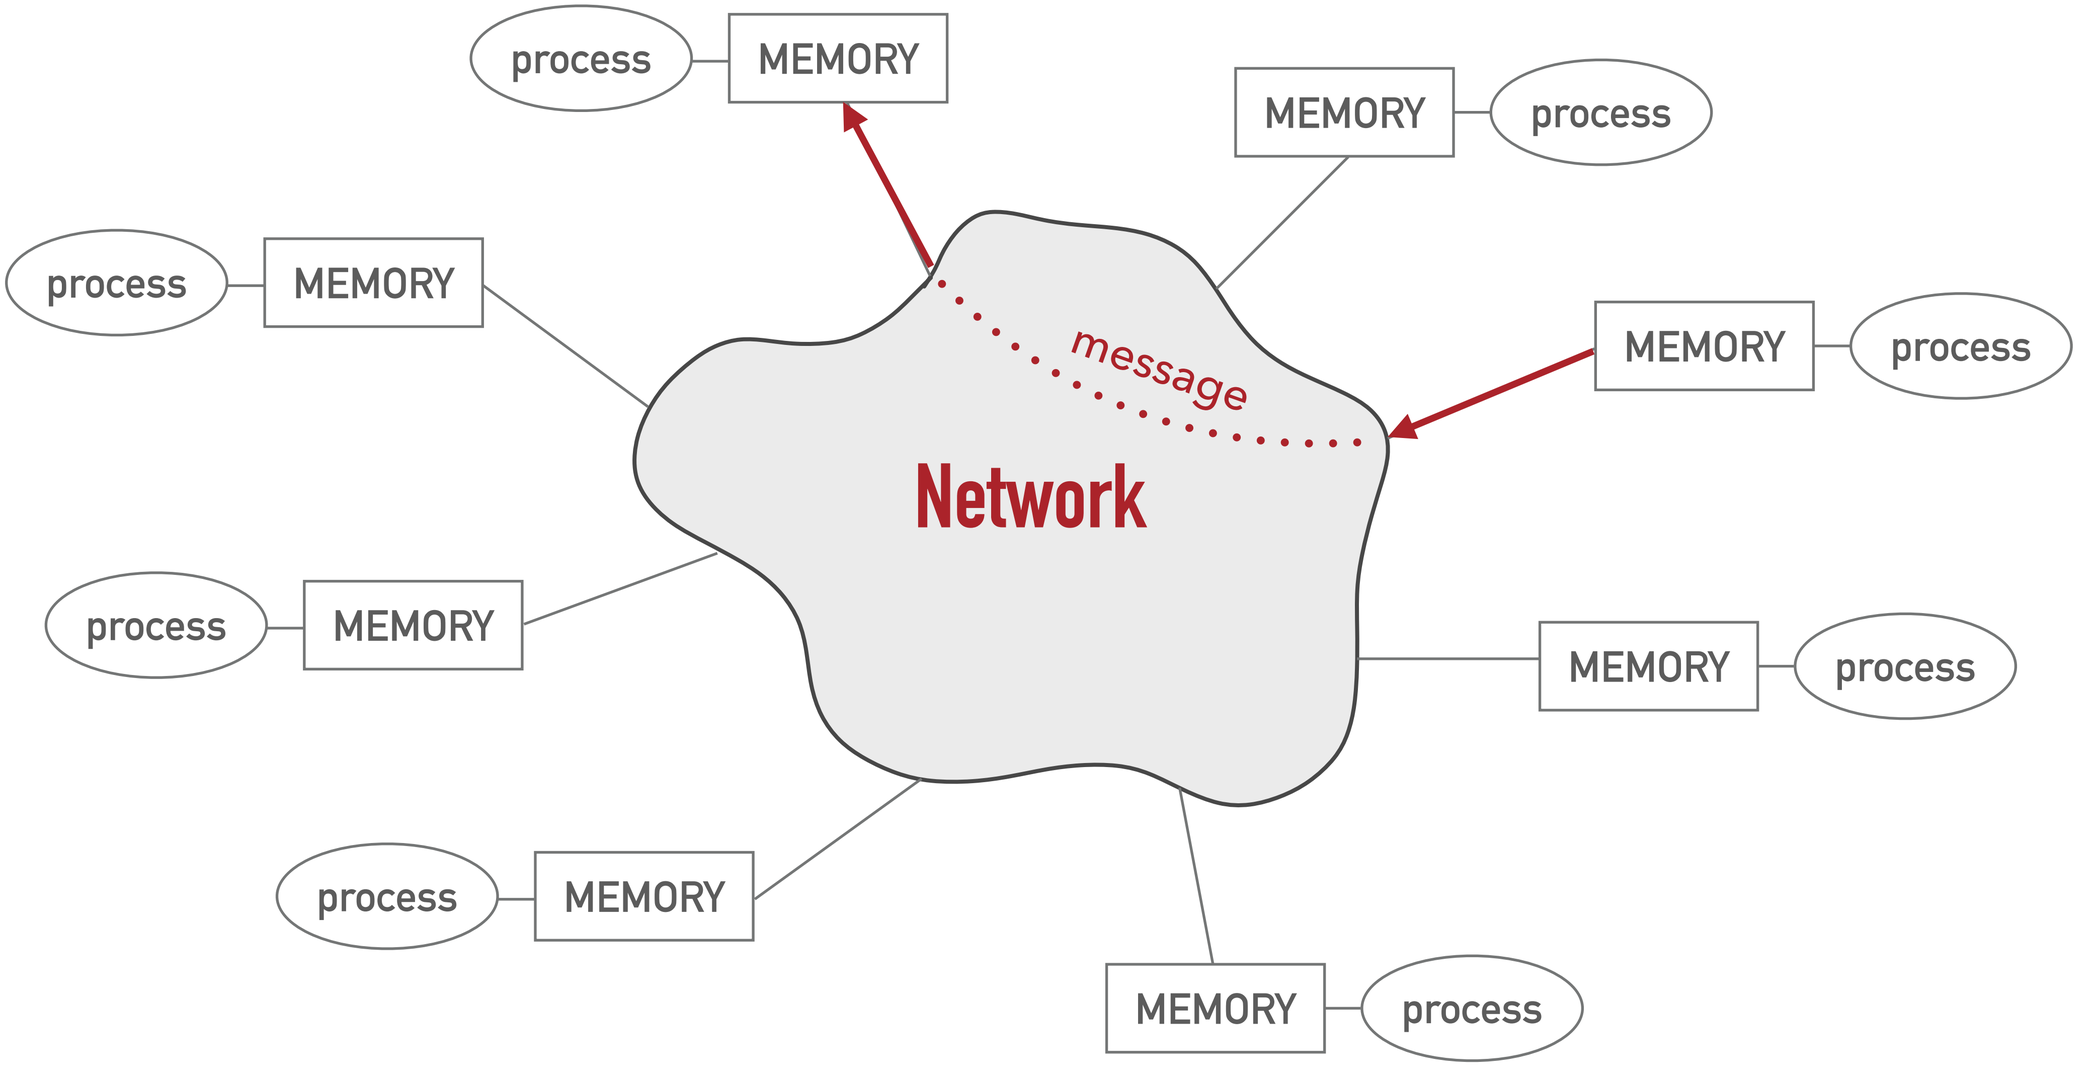

Example of data parallelism - compute the derivative of sin(x)

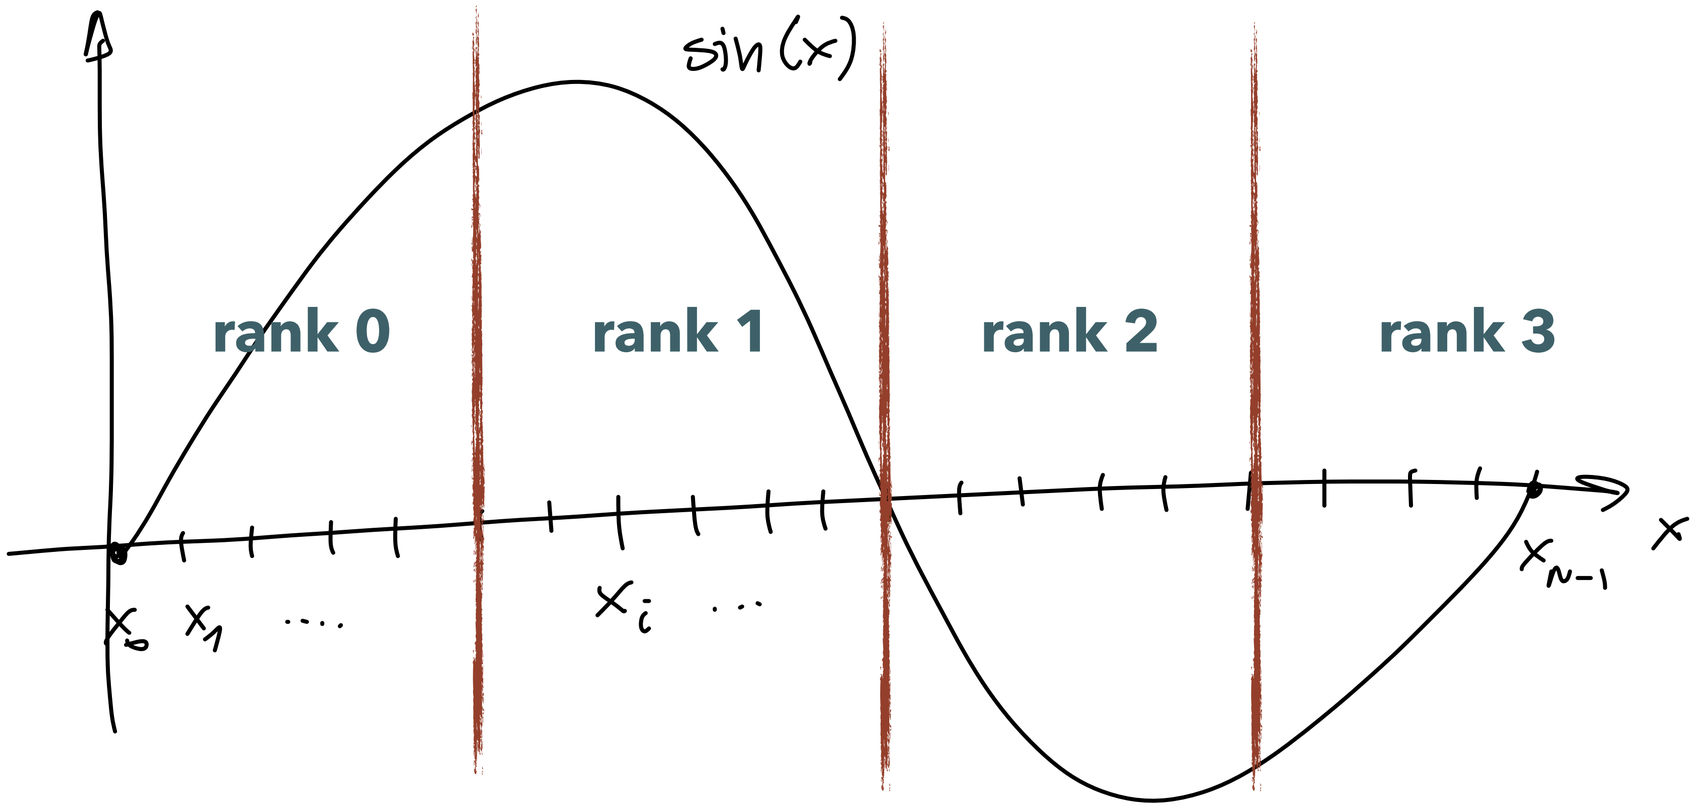

Divide and Conquer: *We can divide the work among available ranks, where each rank performs a fraction of the work*

DERIVATIVE EXAMPLE

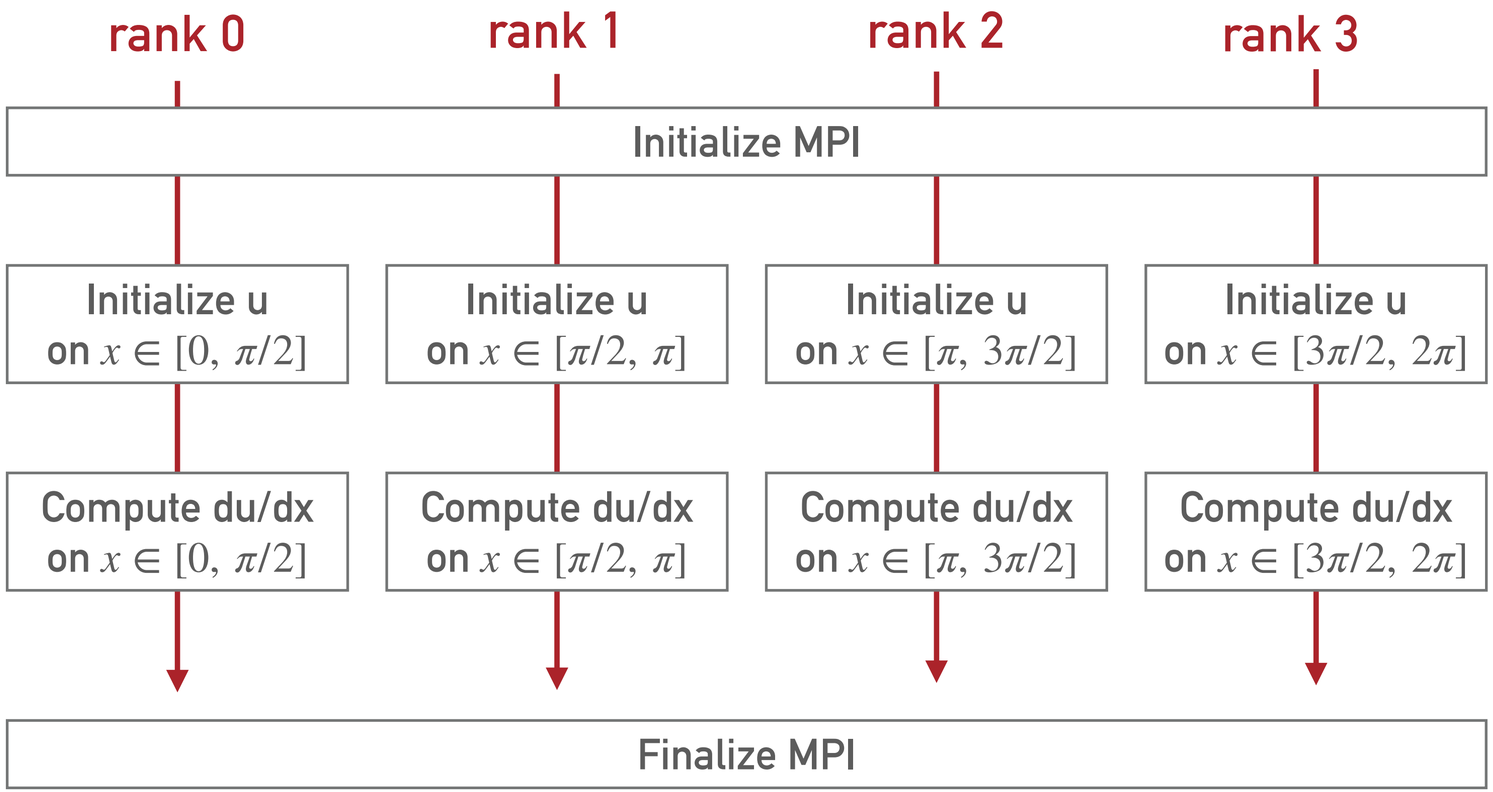

EXAMPLE – INTEGRATE A FUNCTION

Integrate $f(x) = \sin(x)$ *on* $x \in [0,\ 2\pi]$ *using trapezoidal rule:*

$$\int_a^b f(x)\ dx \approx \sum_{i=0}^{N-1} \frac{1}{2} h \left(f(x_i) + f(x_{i+1}) \right)$$

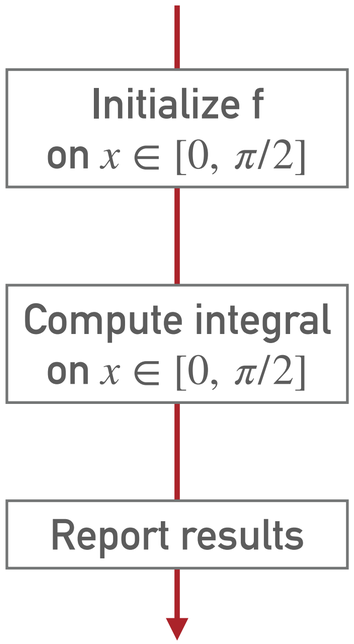

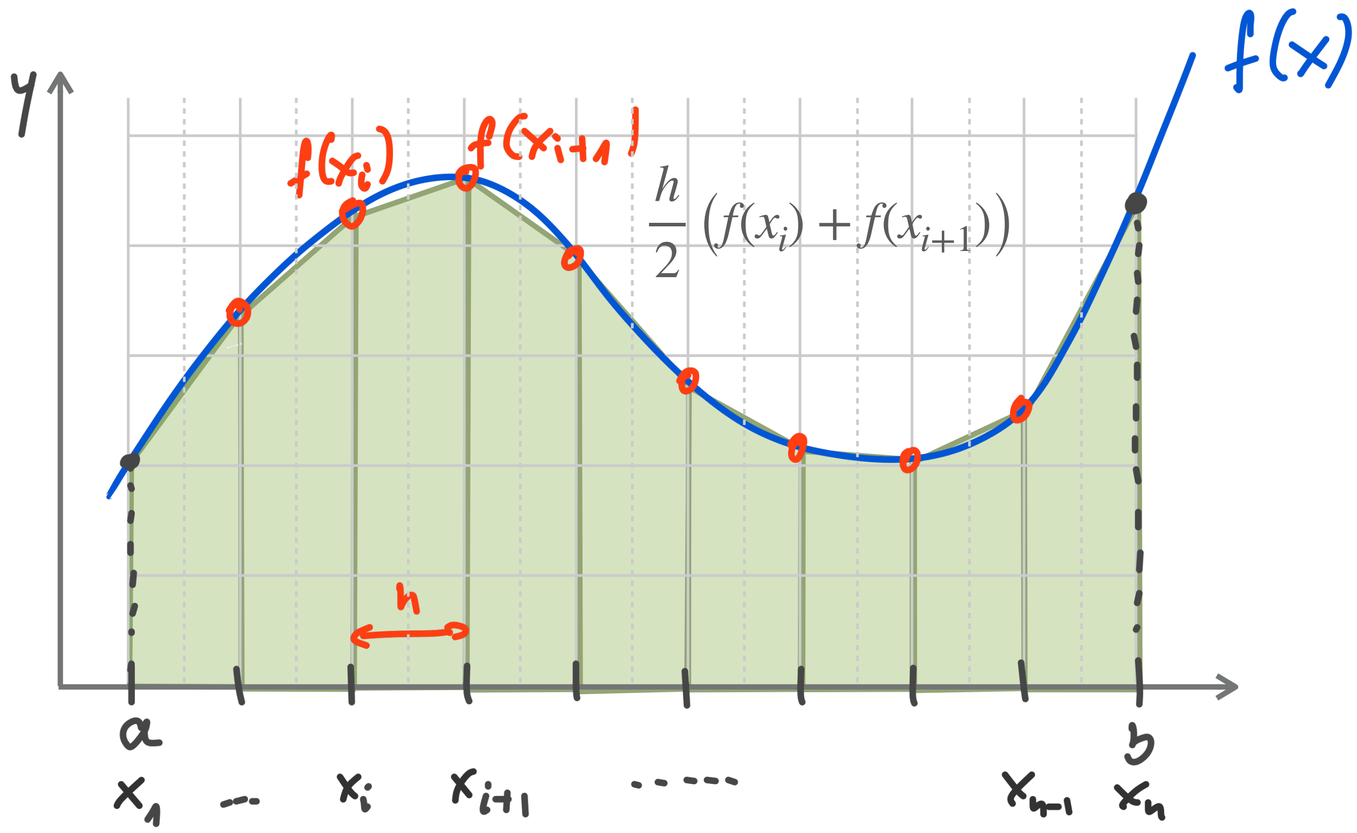

TRAPEZOIDAL RULE EXAMPLE

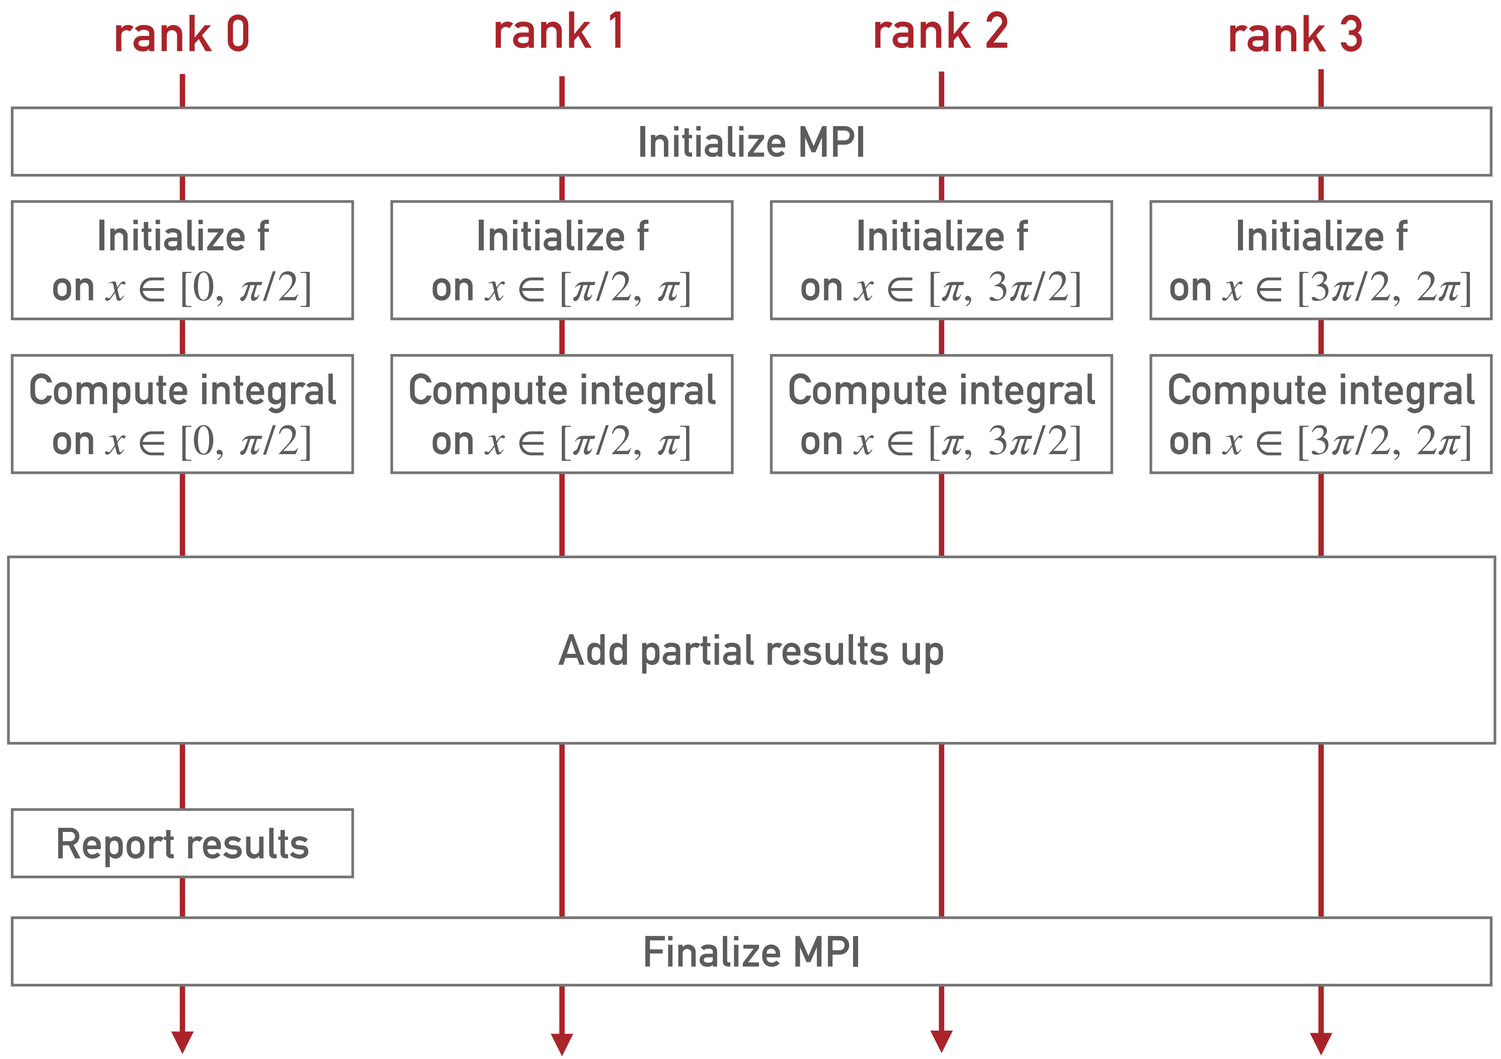

TRAPEZOIDAL RULE EXAMPLE

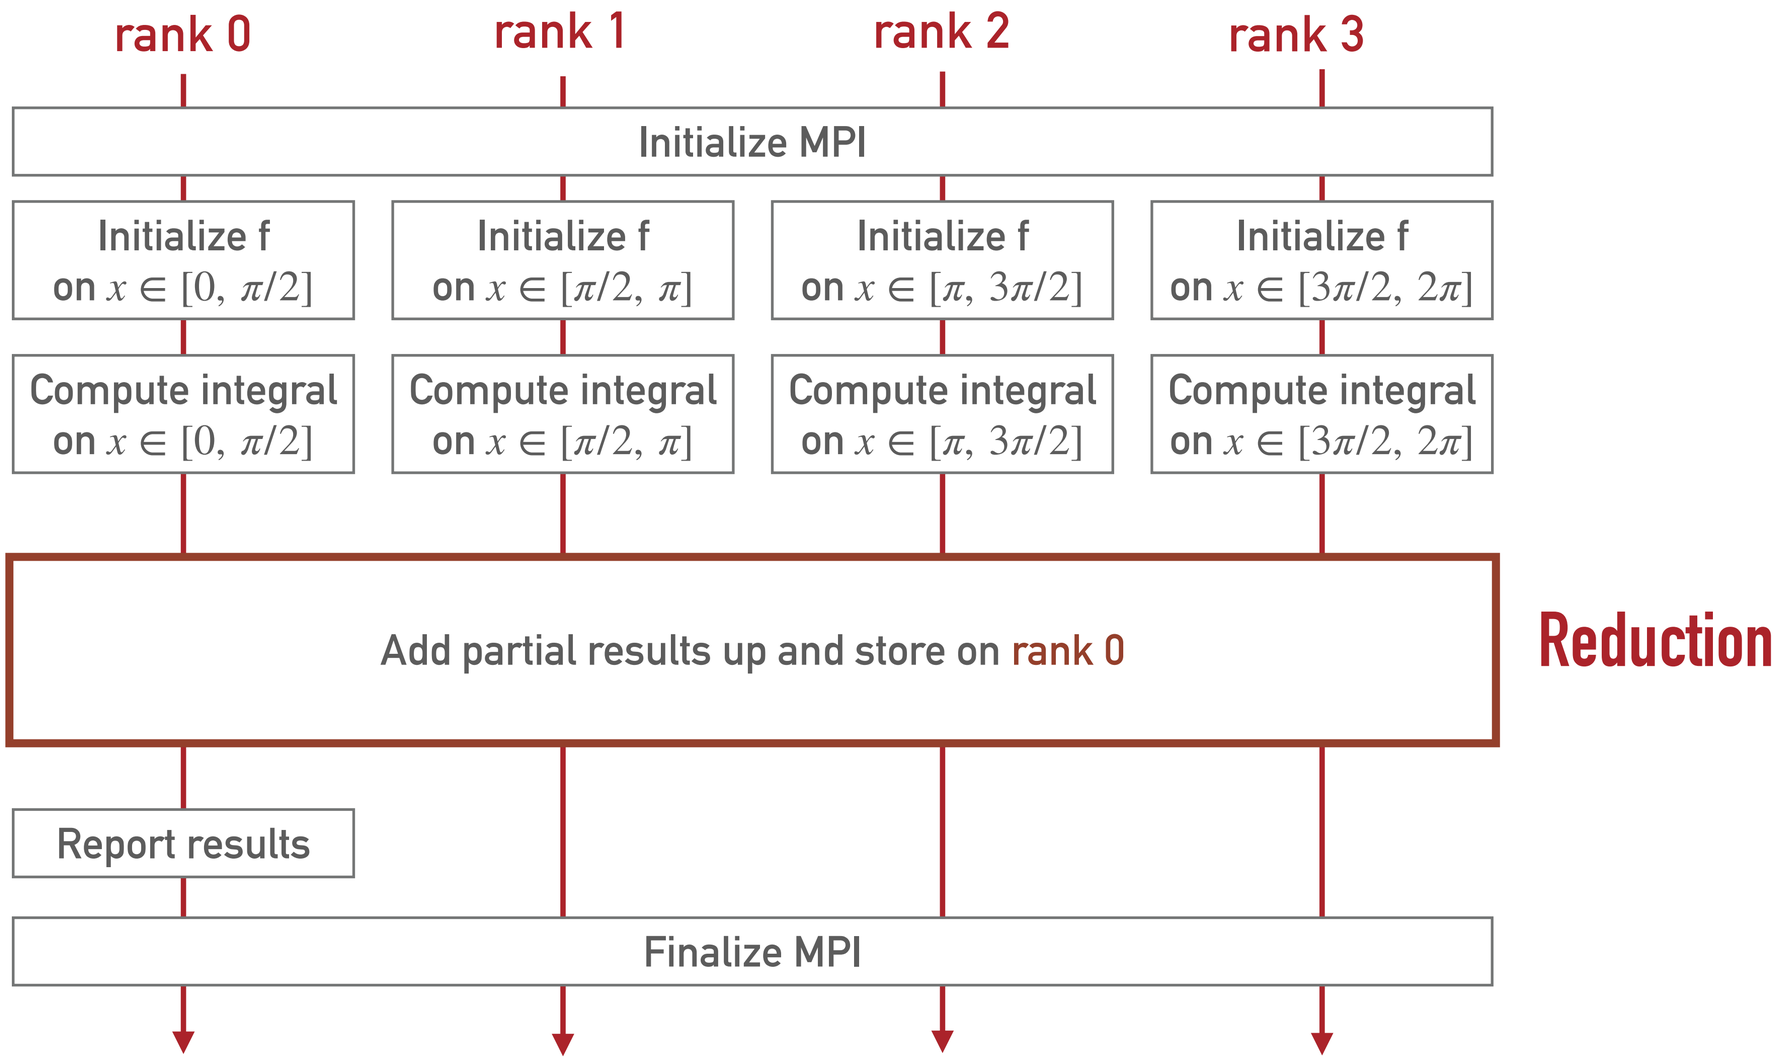

Reduction

COLLECTIVE COMMUNICATION PATTERNS

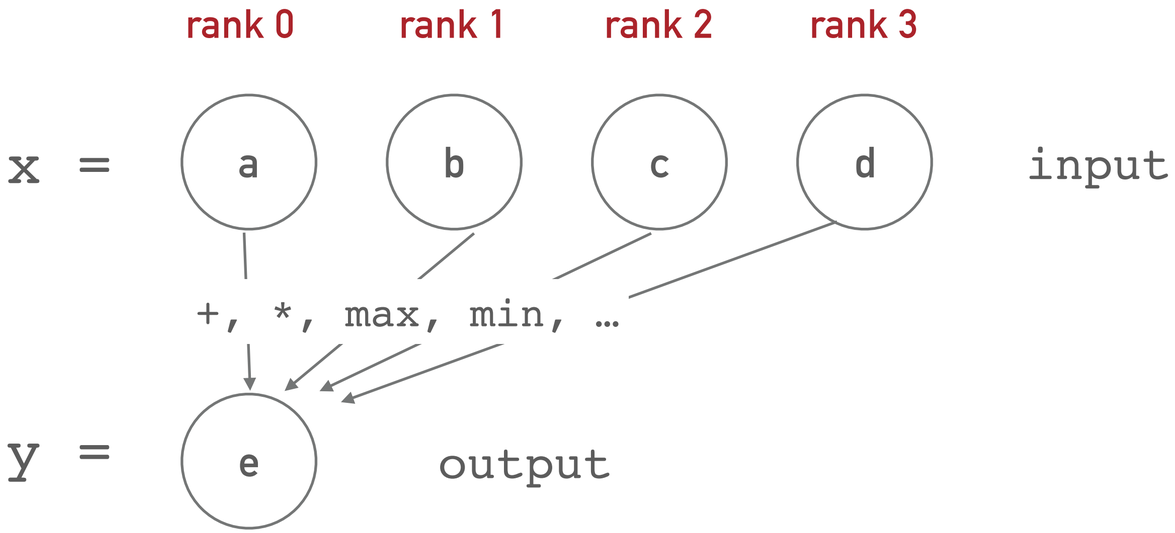

Reduction

```
MPI_Reduce(input, output, count, datatype, operation, root, communicator);
```

which variable to reduce

variable for the result

how many data elements to send per rank

what type?

```
MPI_SUM
MPI_PROD
MPI_MAX
MPI_MIN
…
```

which rank is receiving?

which communicator to use?

TRAPEZOIDAL RULE EXAMPLE

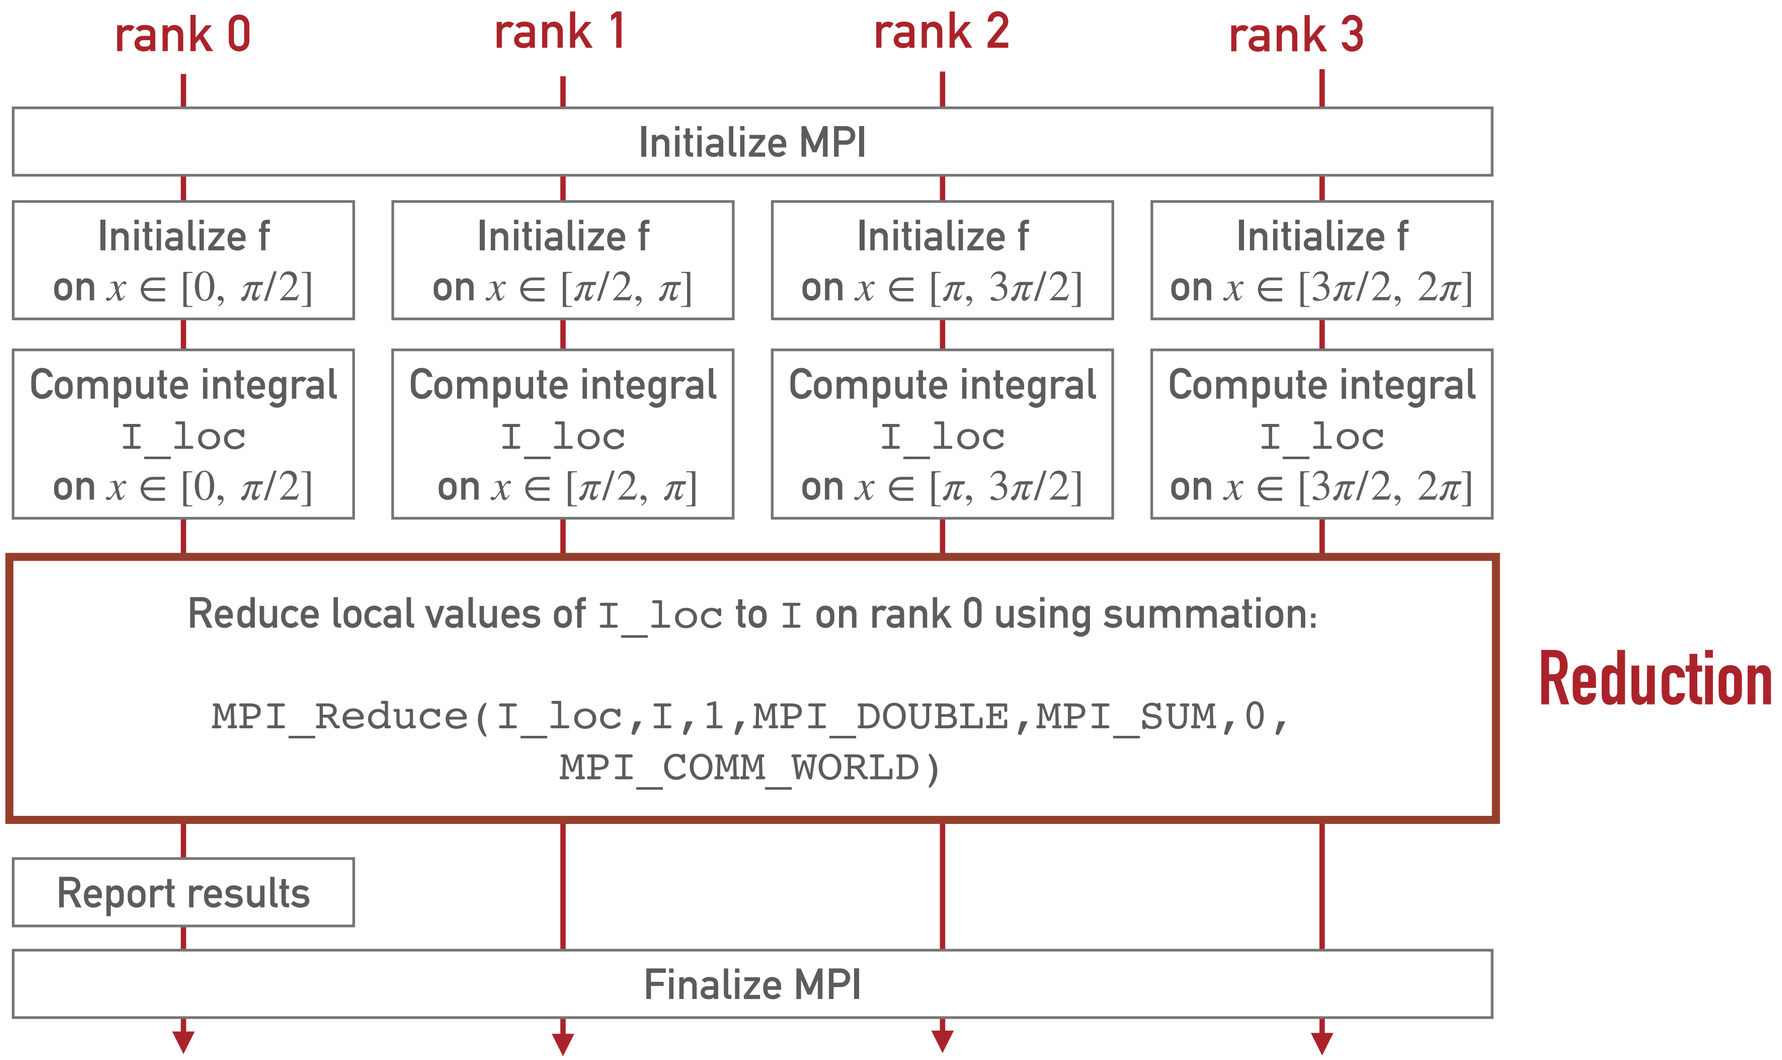

Reduction

COLLECTIVE COMMUNICATION PATTERNS

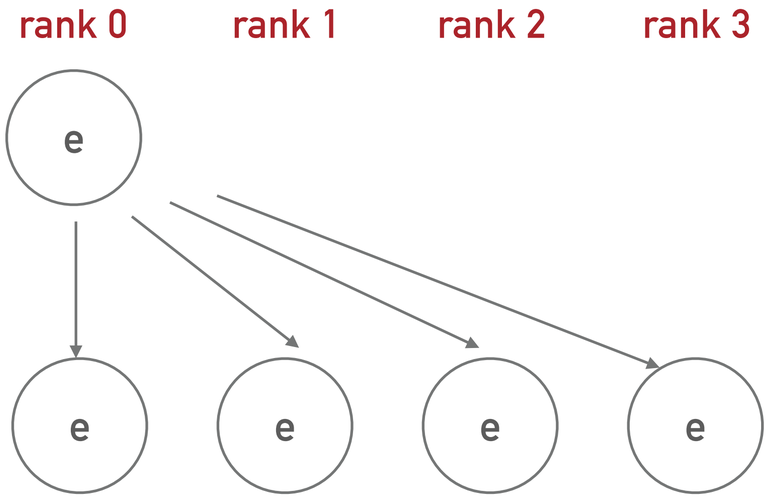

Broadcast

```
MPI_Bcast(data, count, datatype, root, communicator);
```

COLLECTIVE COMMUNICATION PATTERNS

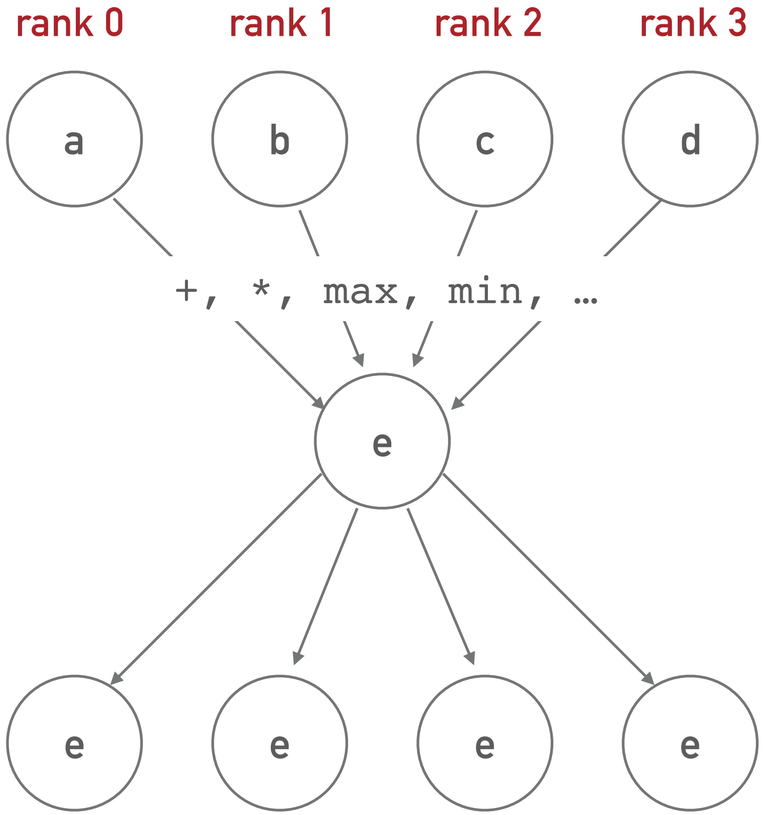

Reduction + Broadcast = Allreduce

```
MPI_Allreduce(input, output, count, datatype, operator, communicator);
```

COLLECTIVE COMMUNICATION PATTERNS

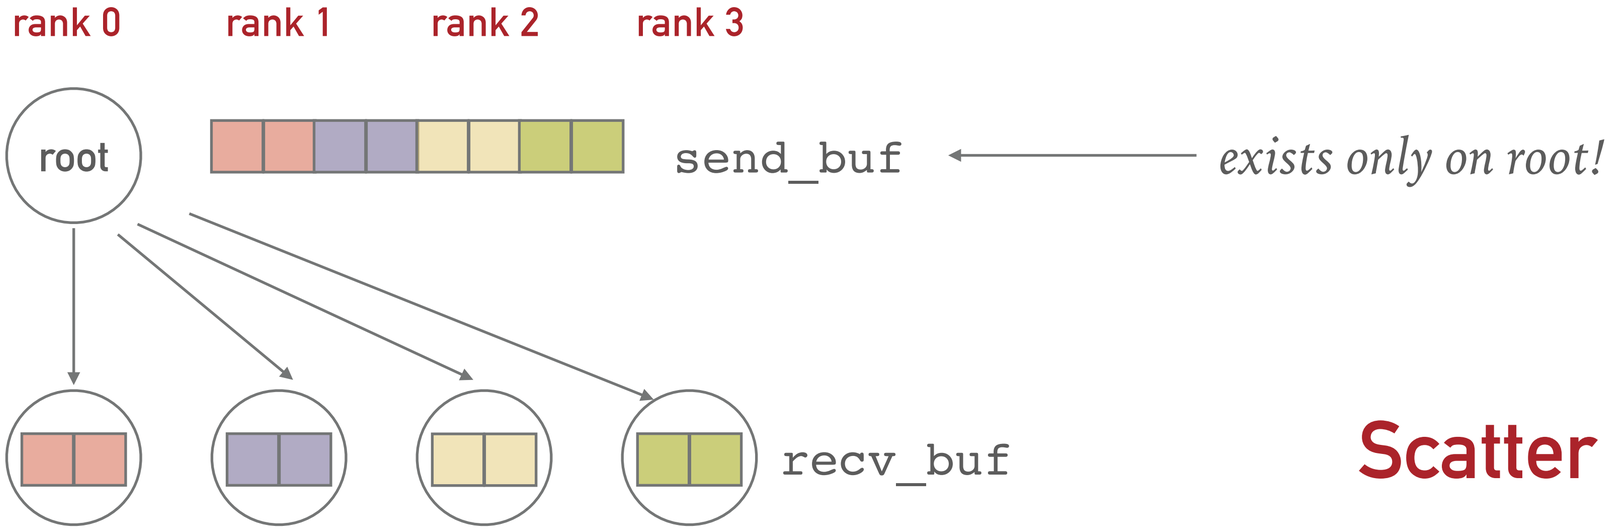

Scatter

```
MPI_Scatter(send_buf, send_count, send_type,…
            recv_buf, recv_count, recv_type,…
            root, communicator);
```

COLLECTIVE COMMUNICATION PATTERNS

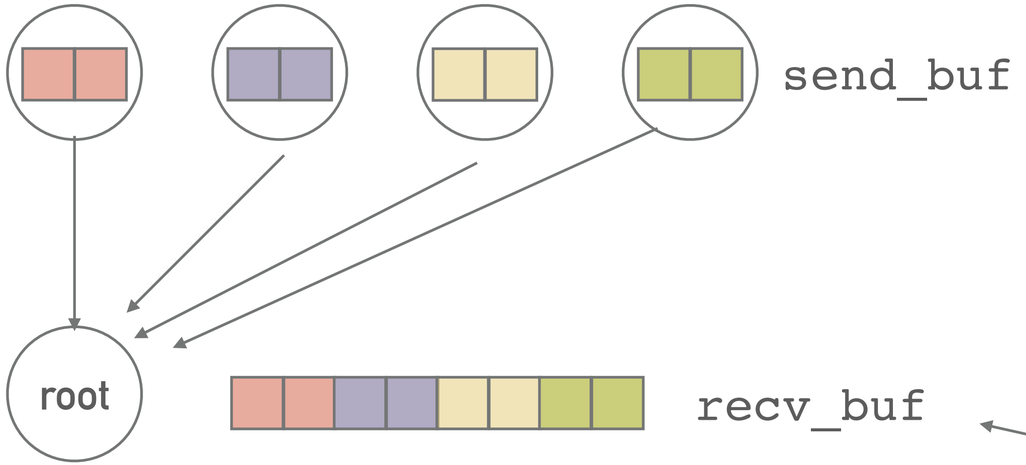

Gather

exists only on root!

```
MPI_Gather(send_buf, send_count, send_type,…
           recv_buf, recv_count, recv_type,…
           root, communicator);
```

COLLECTIVE COMMUNICATION PATTERNS

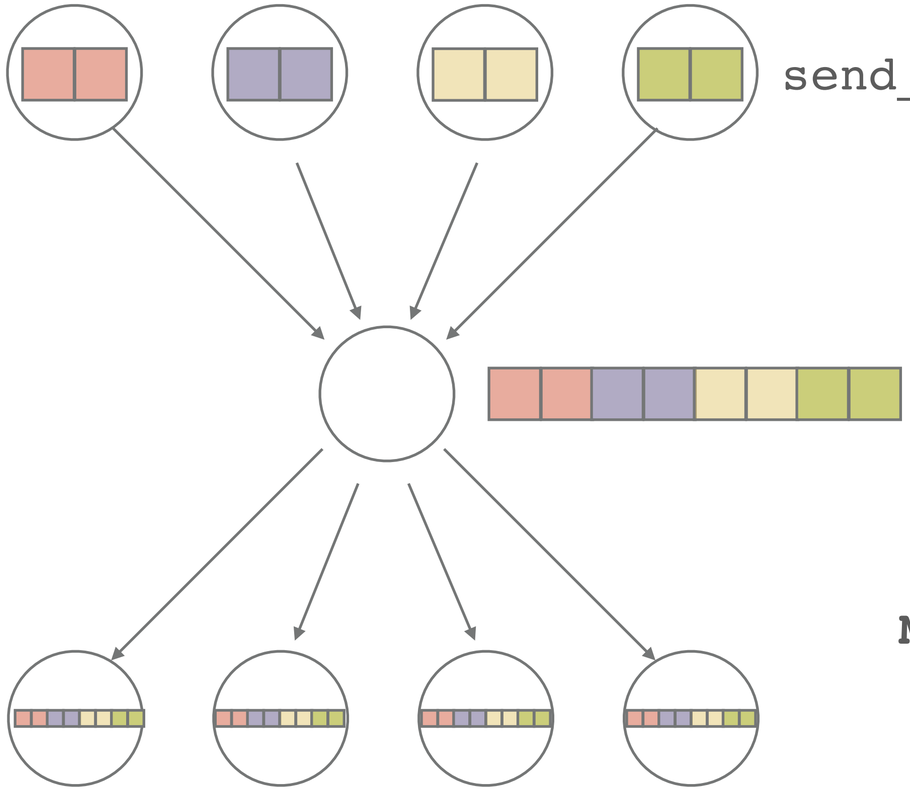

Gather + Scatter = Allgather

```
MPI_Allgather(send_buf, send_count, send_type,…
              recv_buf, recv_count, recv_type,…
              communicator);
```

CLASSROOM EXAMPLE

1. *An instructor is announcing the date of the test.* — broadcast
2. *The students write the test and the instructor is collecting their individual papers.* — gather
3. *The instructor gives the papers back to students.* — scatter
4. *The instructor computes the average score.* — reduction

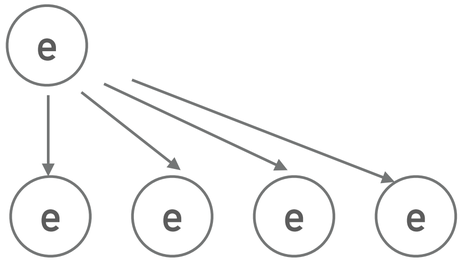

Broadcast

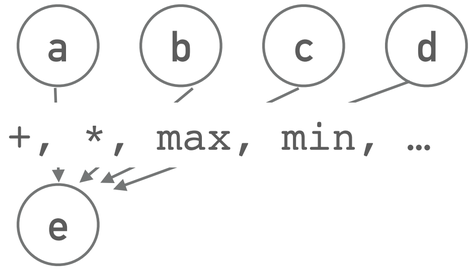

Reduction

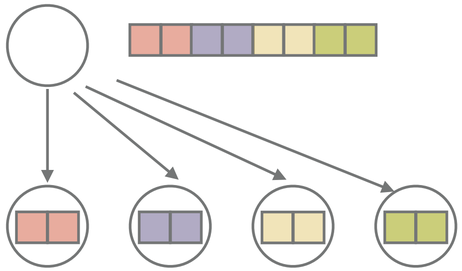

Scatter

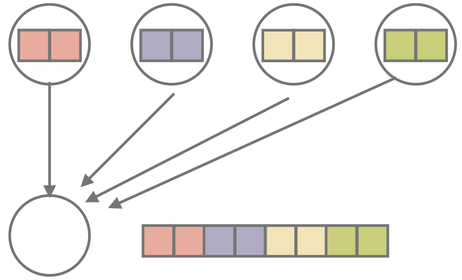

Gather

DERIVATIVE EXAMPLE

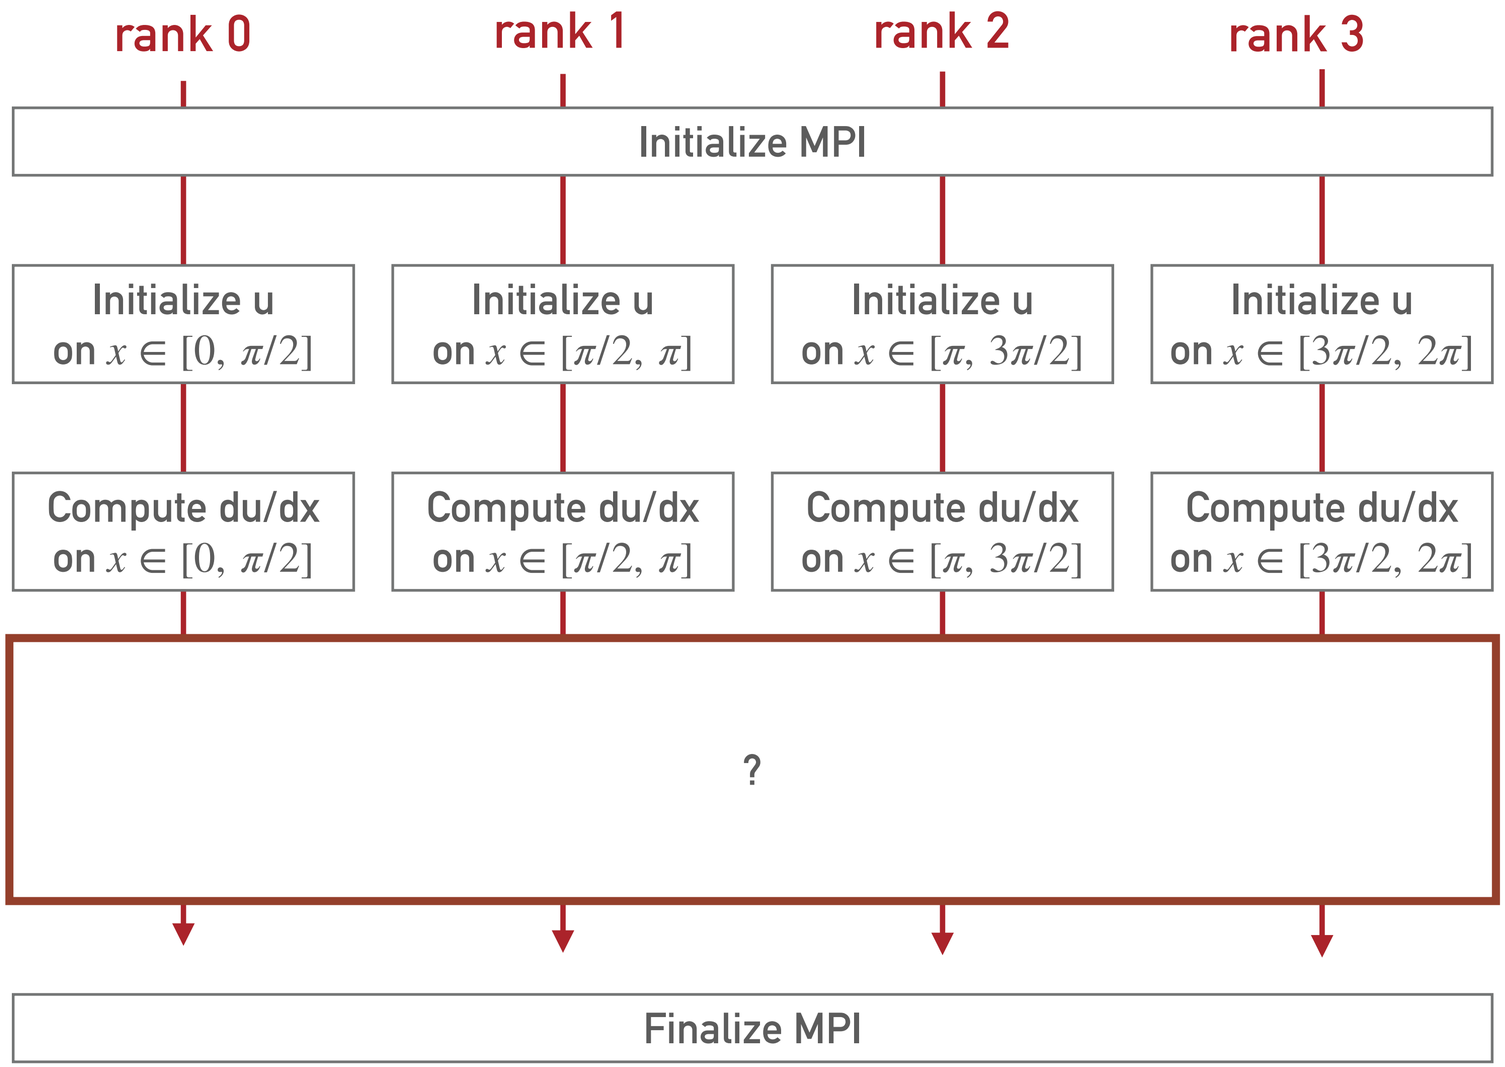

DERIVATIVE EXAMPLE

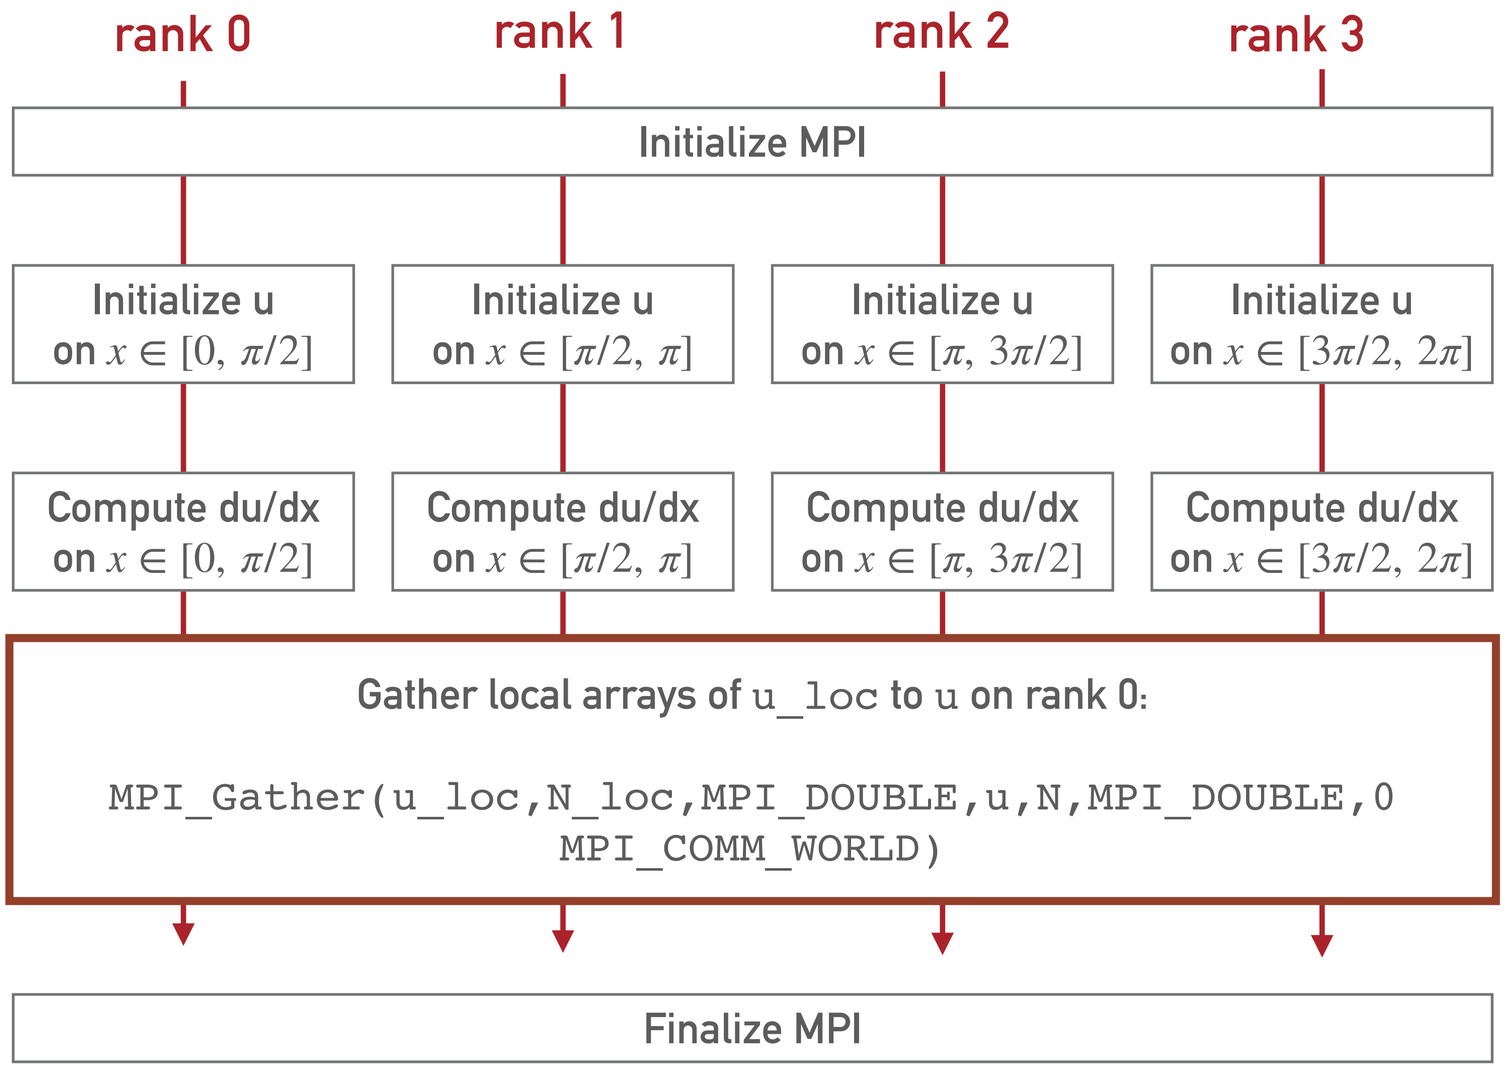

EXAMPLE

Imagine you need to compute a standard deviation of a large set of numbers saved in a file. How would you design a parallel algorithm using Reduce, Broadcast, Scatter, Gather, … functionalities?

$$\sigma = \sqrt{\frac{1}{N}\sum_{i=0}^{N-1}(x_i - \mu)} \qquad \textit{where} \qquad \mu = \frac{1}{N}\sum_{i=0}^{N-1} x_i$$